Learning Word Embeddings

MMCi Block 5

Matthew Engelhard

From sentence to sequence of vectors.

- Look up words individually to obtain their vectors
- But where do the vectors come from?

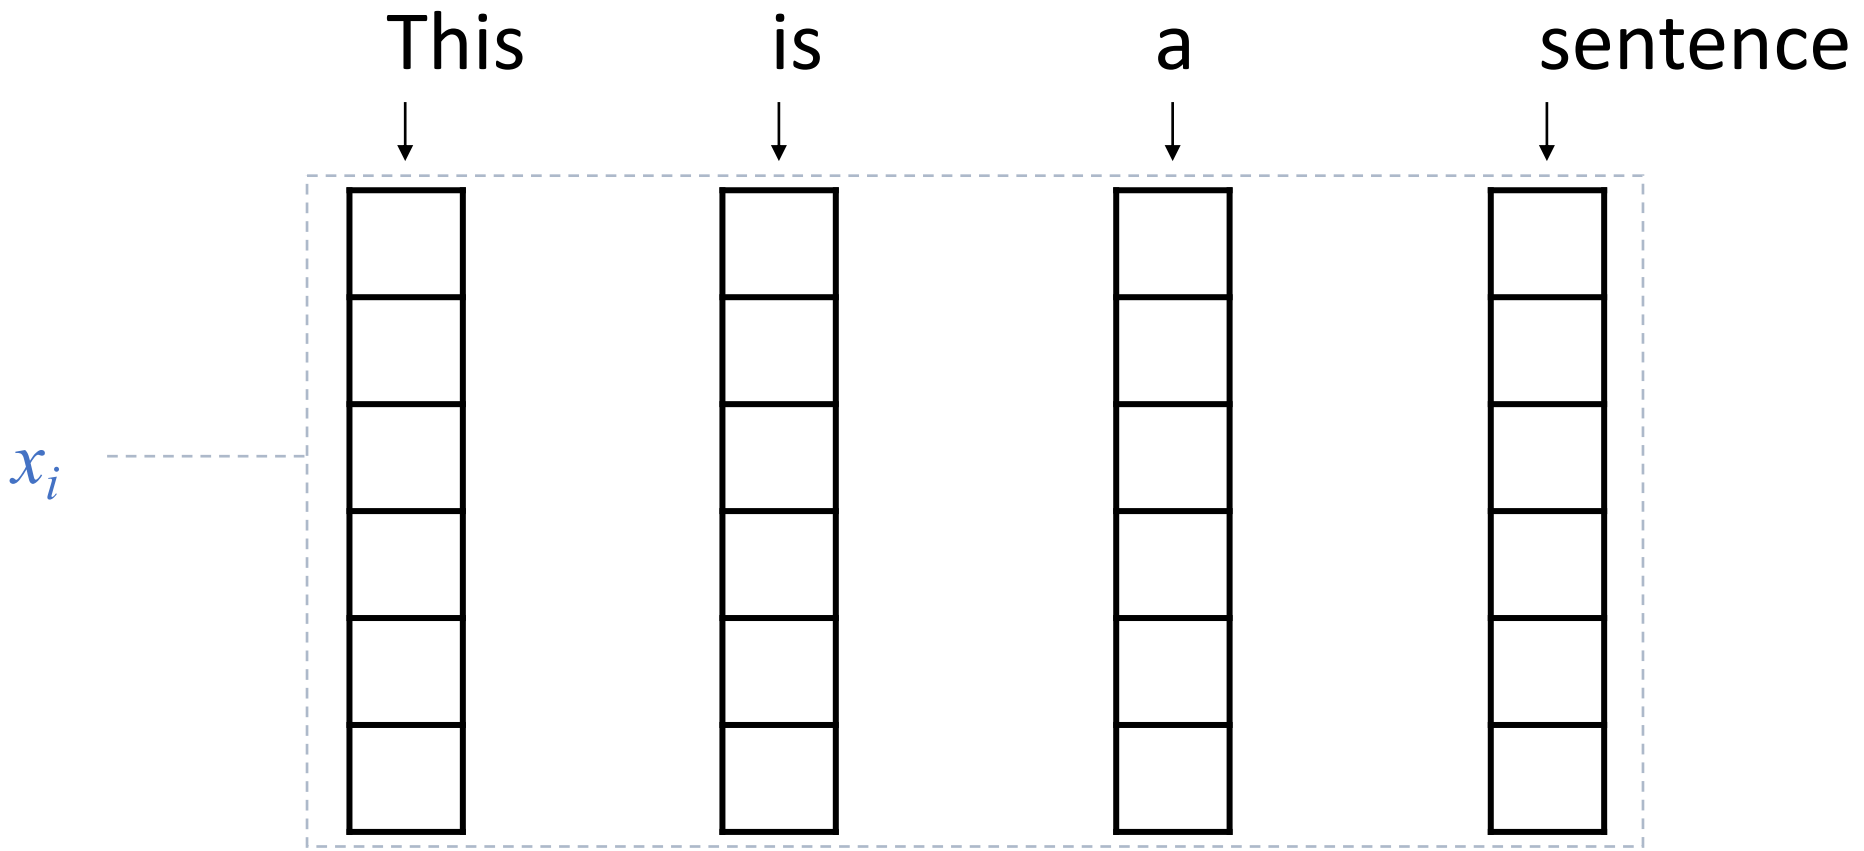

Visualizing Word Embeddings

If the meaning is similar, the vectors (i.e. locations) should be similar!

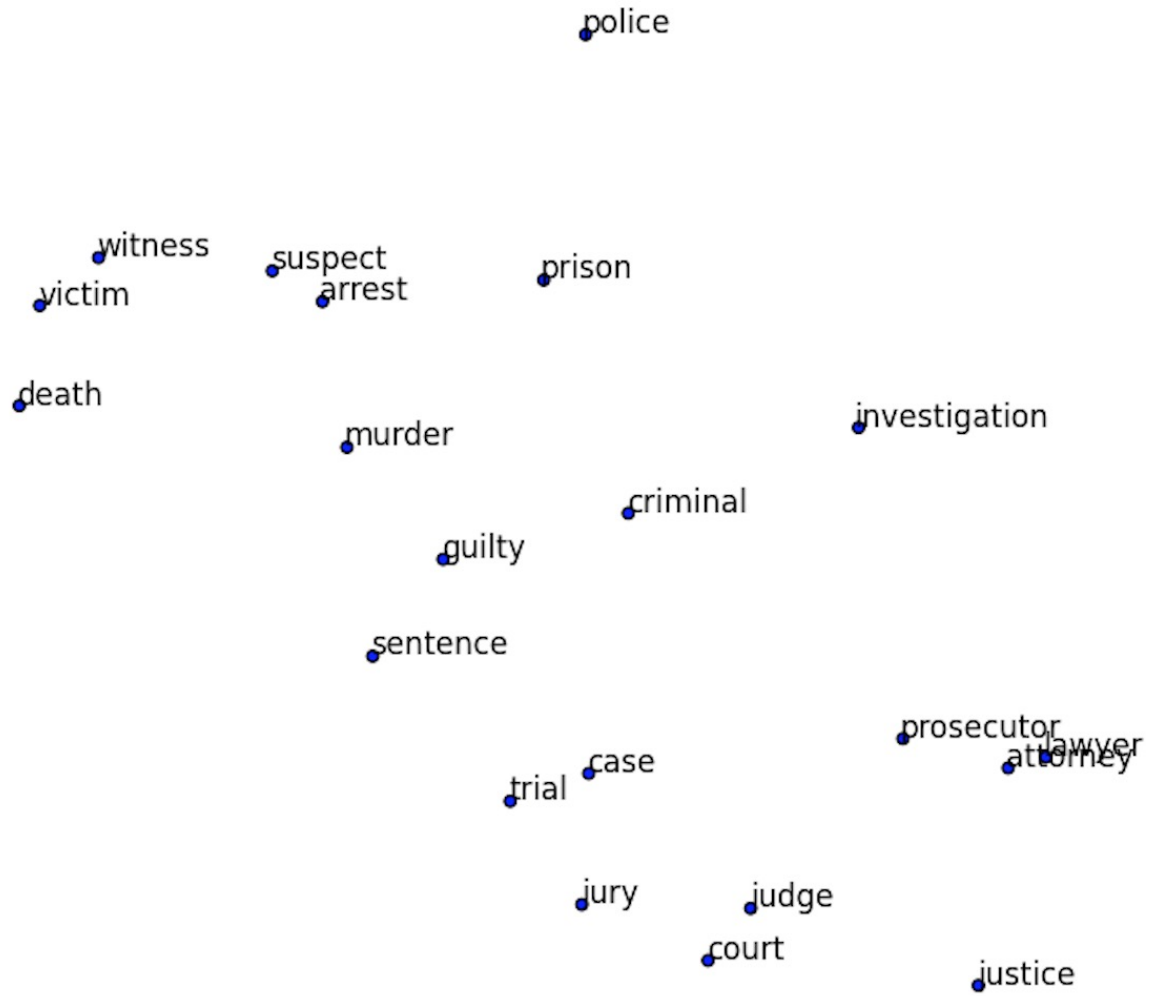

How are word embeddings learned?

KEY IDEA: words are *defined* by the context in which they appear

A man strolls down the street

A woman strolls down the street

A child strolls down the street

A crocodile strolls down the street

A banana strolls down the street

A concept strolls down the street

KEY IDEA: words are *defined* by the context in which they appear

-> if words are always exchangeable, they must have very similar meaning

learn word meaning like an adult:
explicit definitions

https://www.parenting.com/activities/baby/teach-baby-to-talk/

learn word meaning like an child:
implicit definitions from context

So how do we learn spatial locations for each word?

We want: a vector for each word that allows us to predict its context (i.e. what other words are likely/unlikely to be around it)

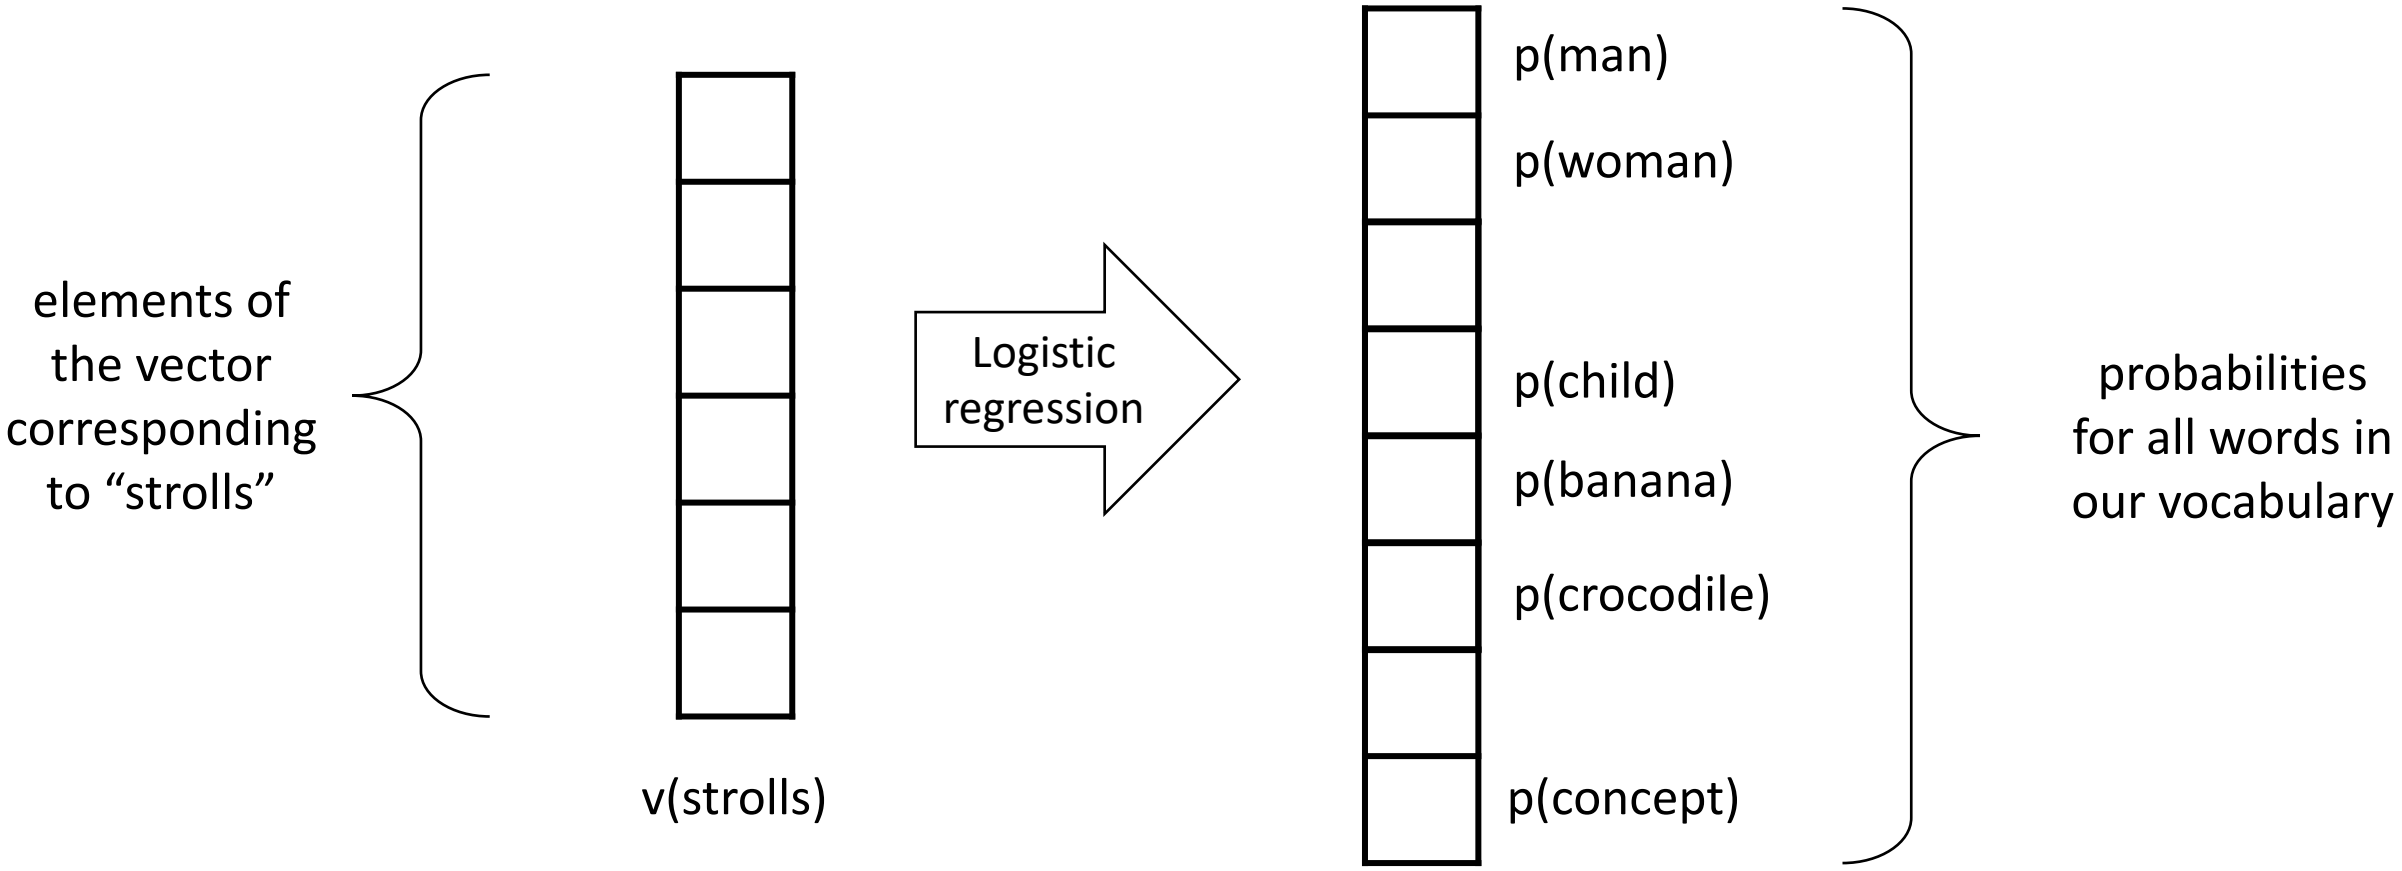

Predict Context Words from Input Words

{input word, context word}

{strolls, man}
{strolls, woman}
{swims, crocodile}
{swims, fish}
{flies, bird}
{flies, plane}

We define a context word as one that appears inside a fixed-length window around the input word in our training corpus.

(e.g. Wikipedia)

A man strolls down the street.

input context

So how do we learn spatial locations for each word?

We want: a vector for each word that allows us to predict its context (i.e. what other words are likely/unlikely to be around it)

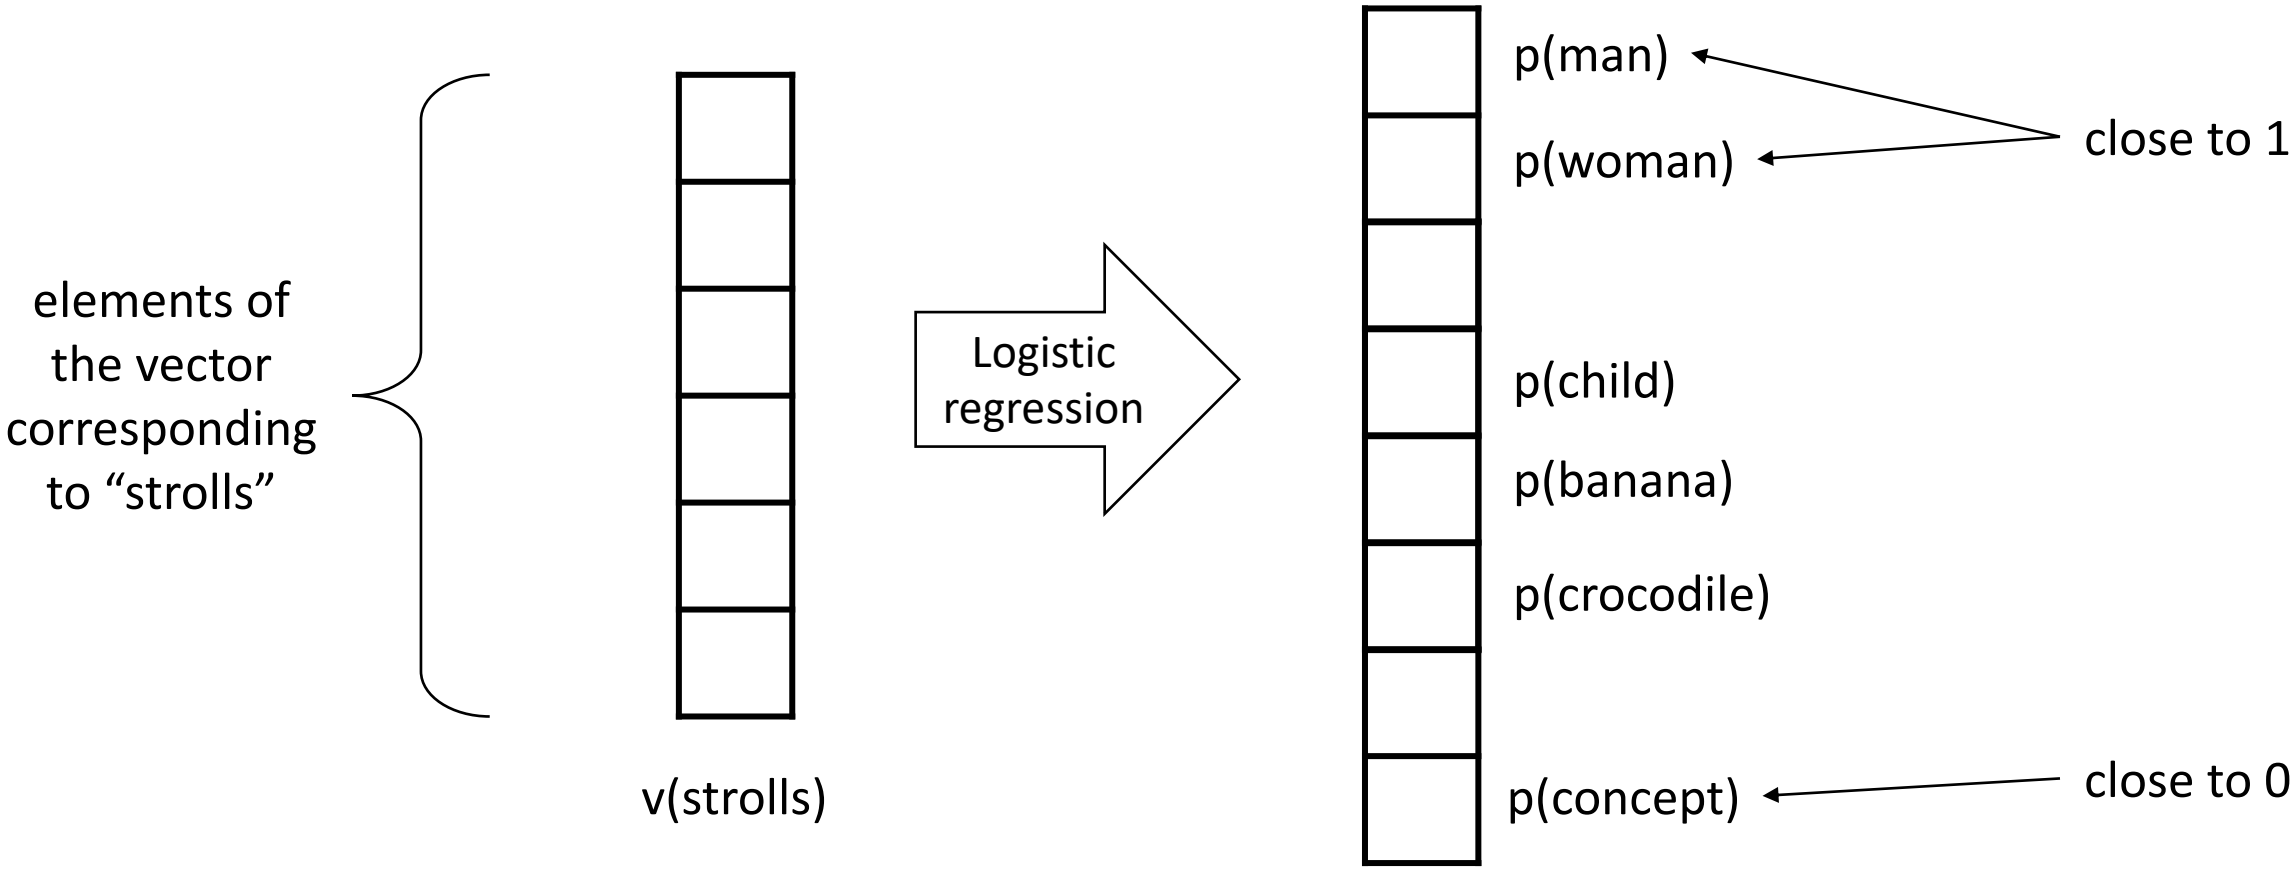

So how do we learn spatial locations for each word?

We want: a vector for each word that allows us to predict its context (i.e. what other words are likely/unlikely to be around it)

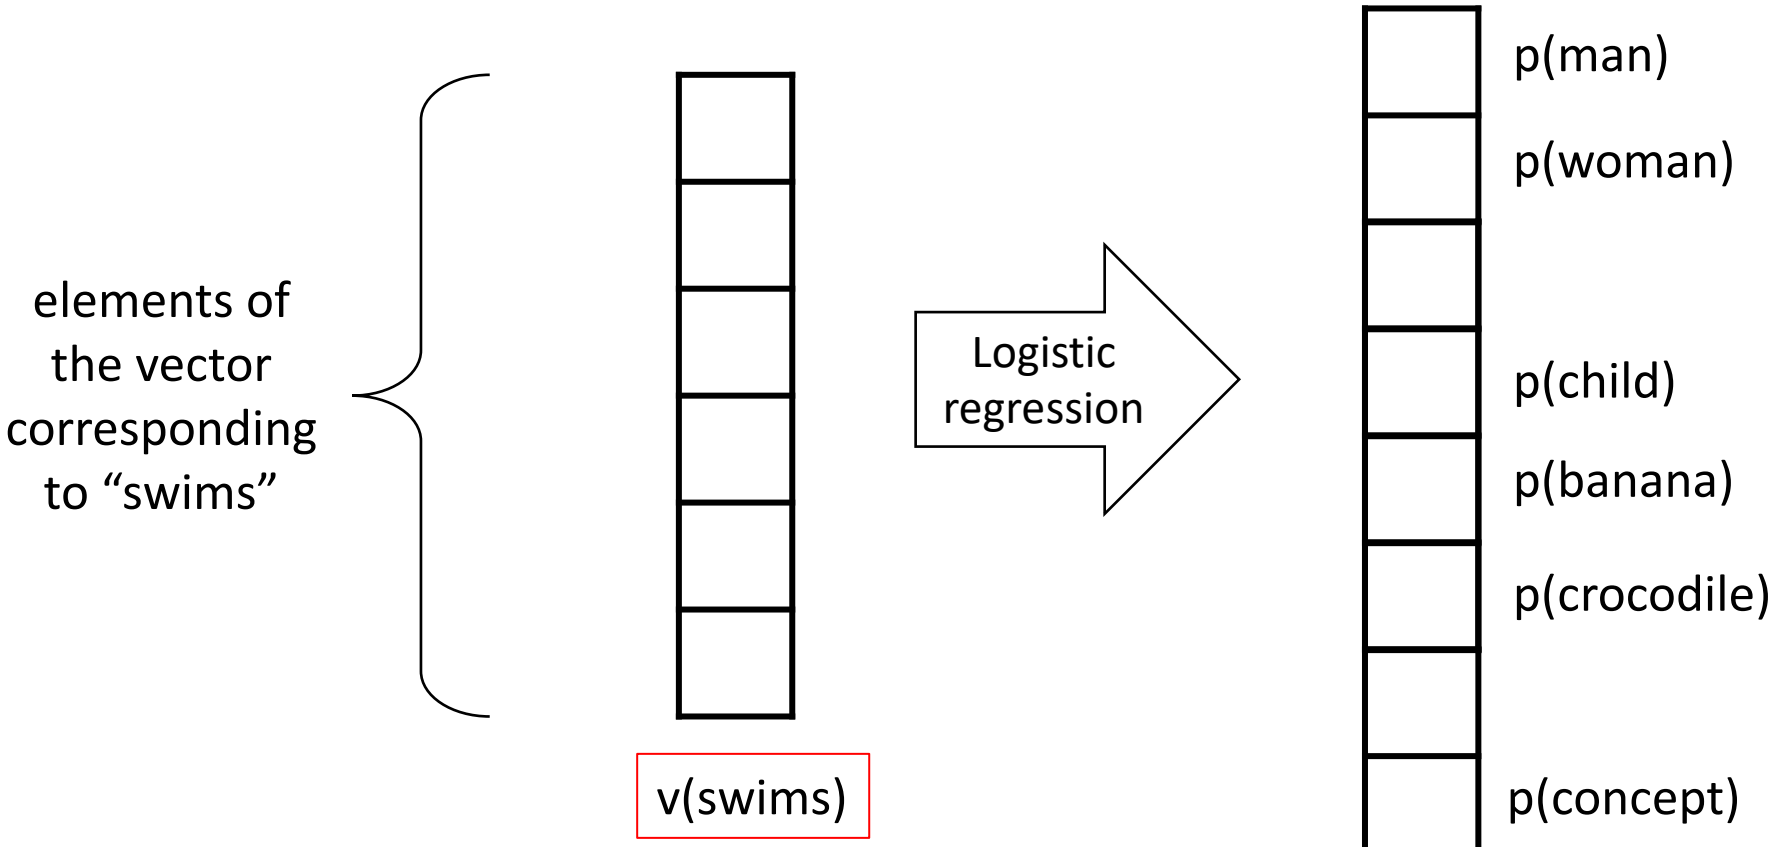

So how do we learn spatial locations for each word?

We want: a vector for each word that allows us to predict its context (i.e. what other words are likely/unlikely to be around it)

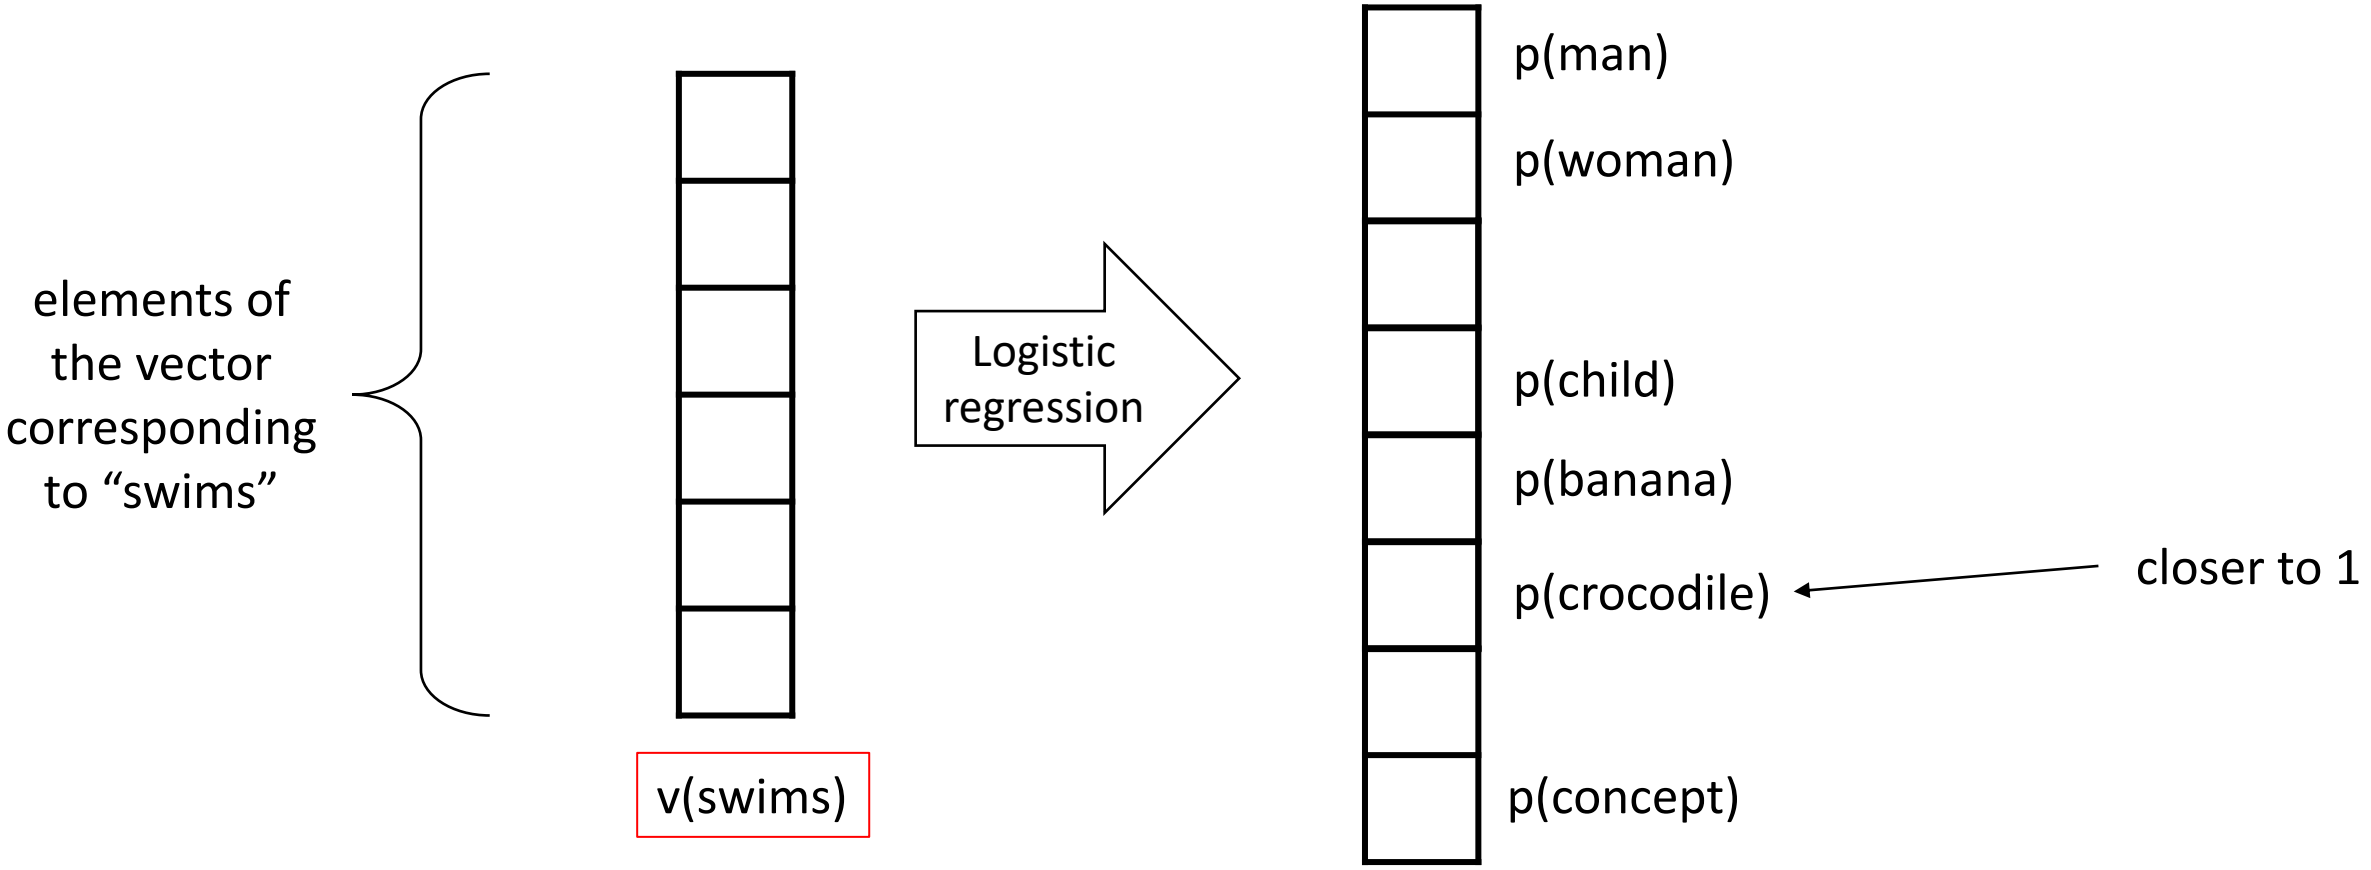

So how do we learn spatial locations for each word?

We want: a vector for each word that allows us to predict its context (i.e. what other words are likely/unlikely to be around it)

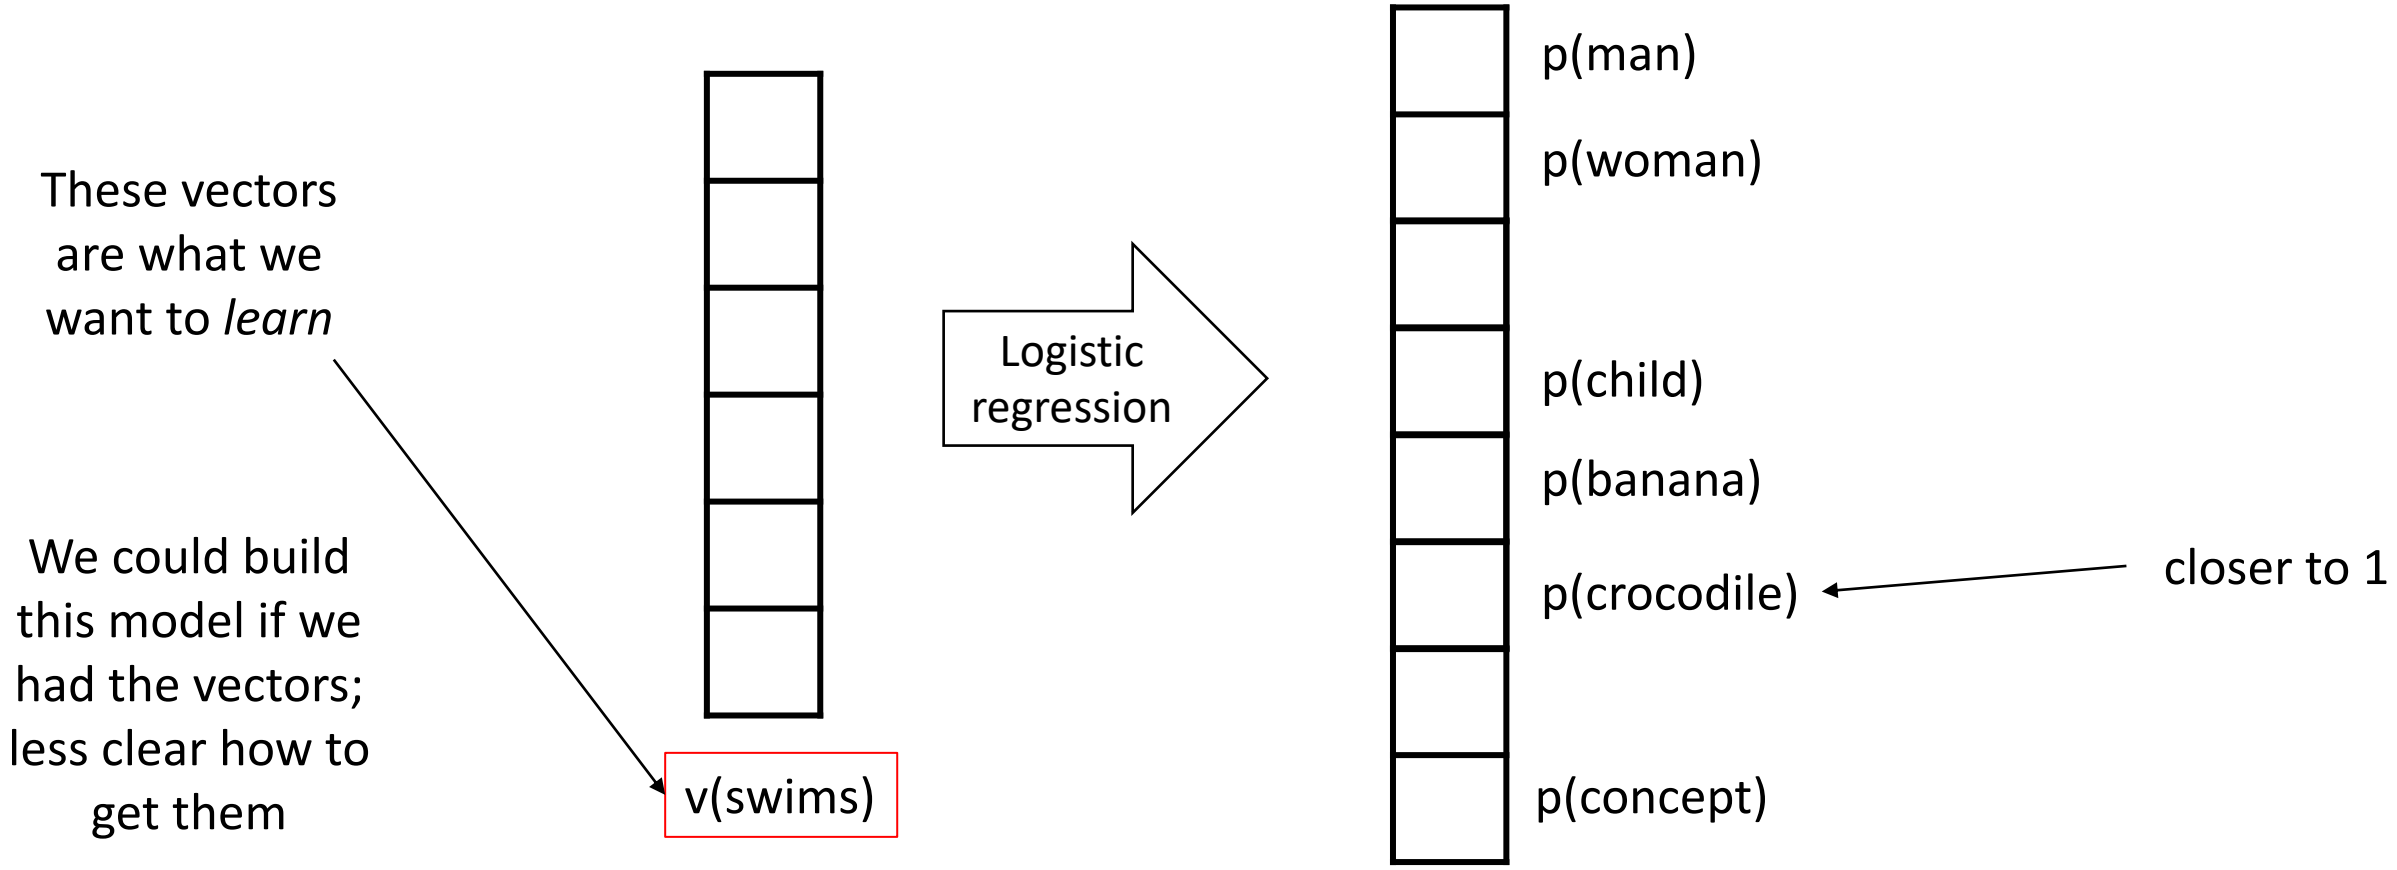

Recall: Multi-Class Logistic Regression

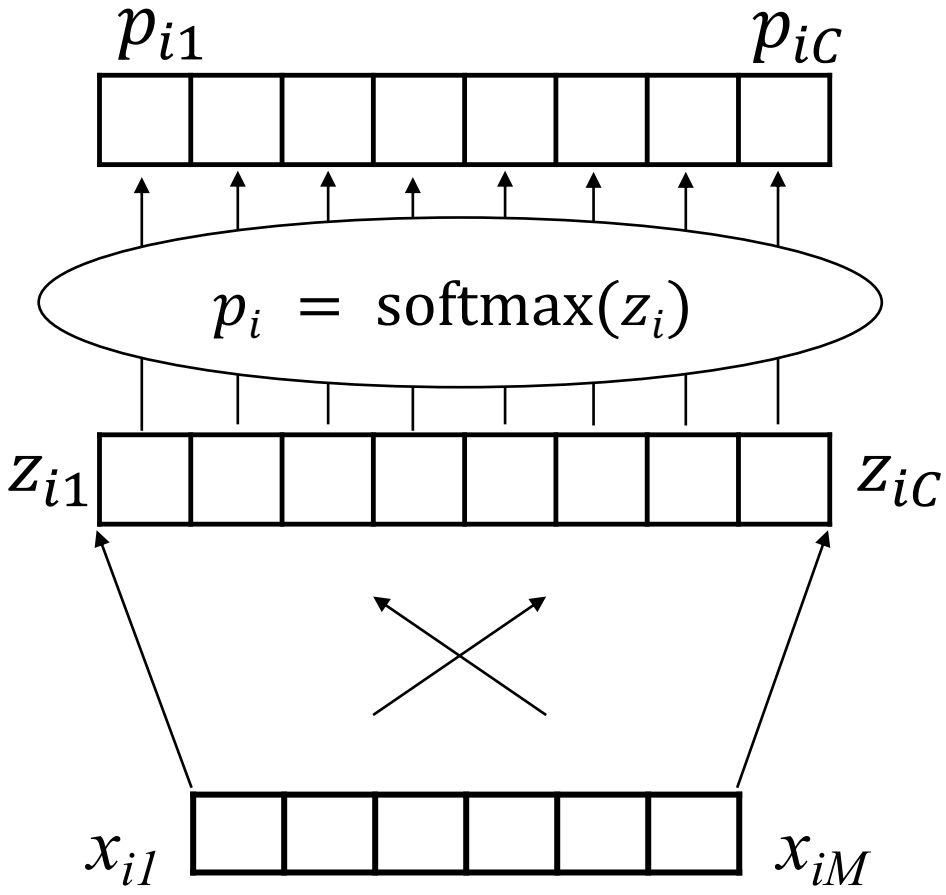

$$p_{ij} = \frac{e^{z_{ij}}}{\sum_{c=1}^{C} e^{z_{ic}}}$$

Recall: Multi-Class Logistic Regression

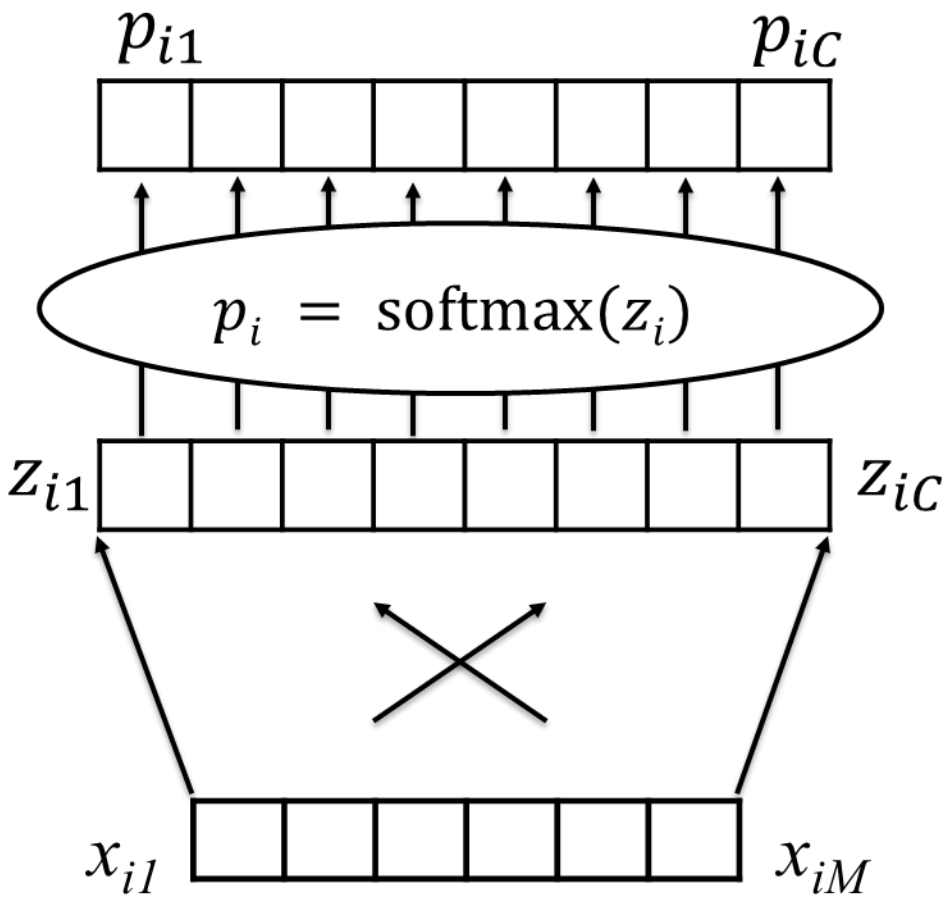

Recall: Multi-Class Logistic Regression

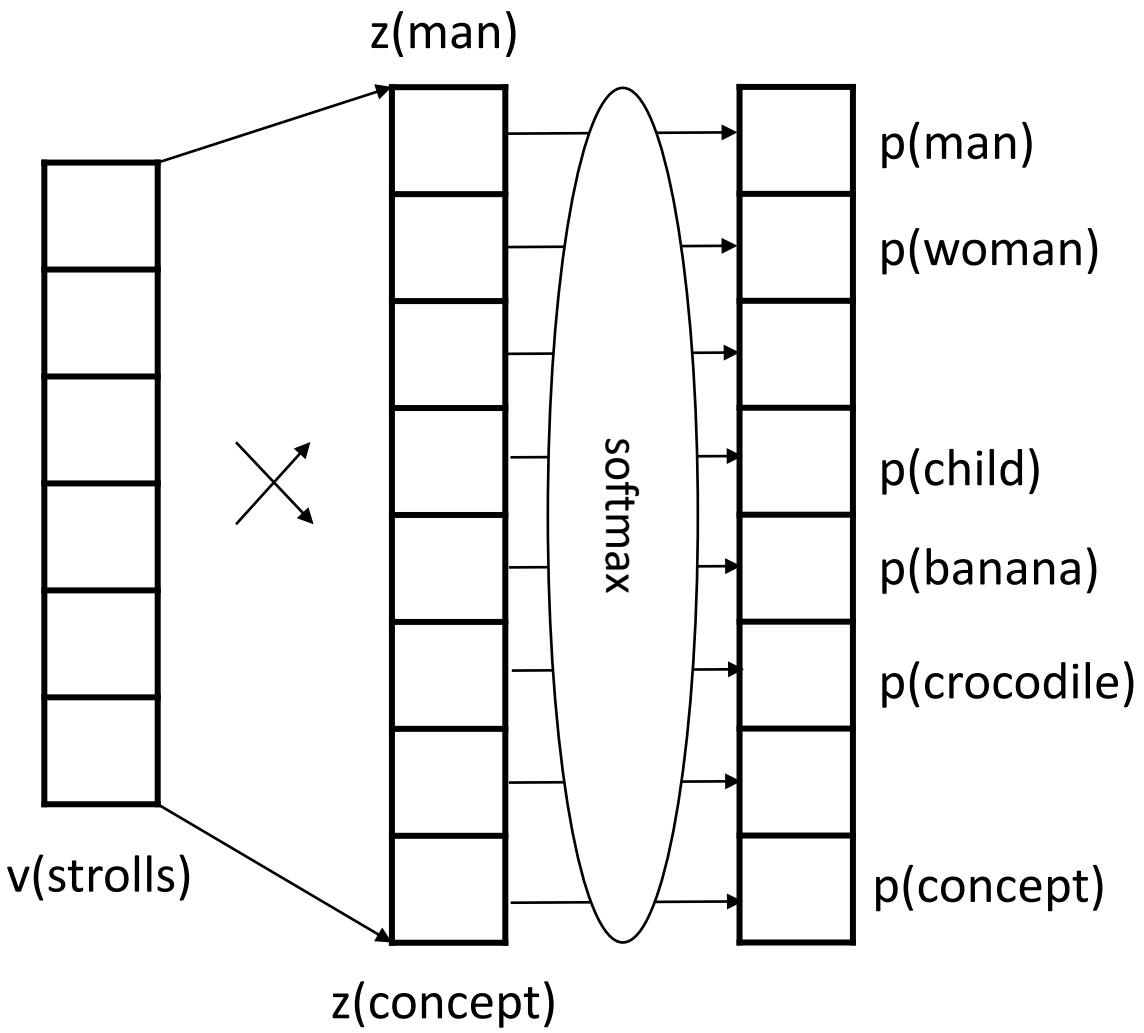

We want: word vectors that allow us to predict their likely context

But again, how do we *learn* these vectors?

Let's take a step back: we'll focus on understanding how we can predict context words based on input words

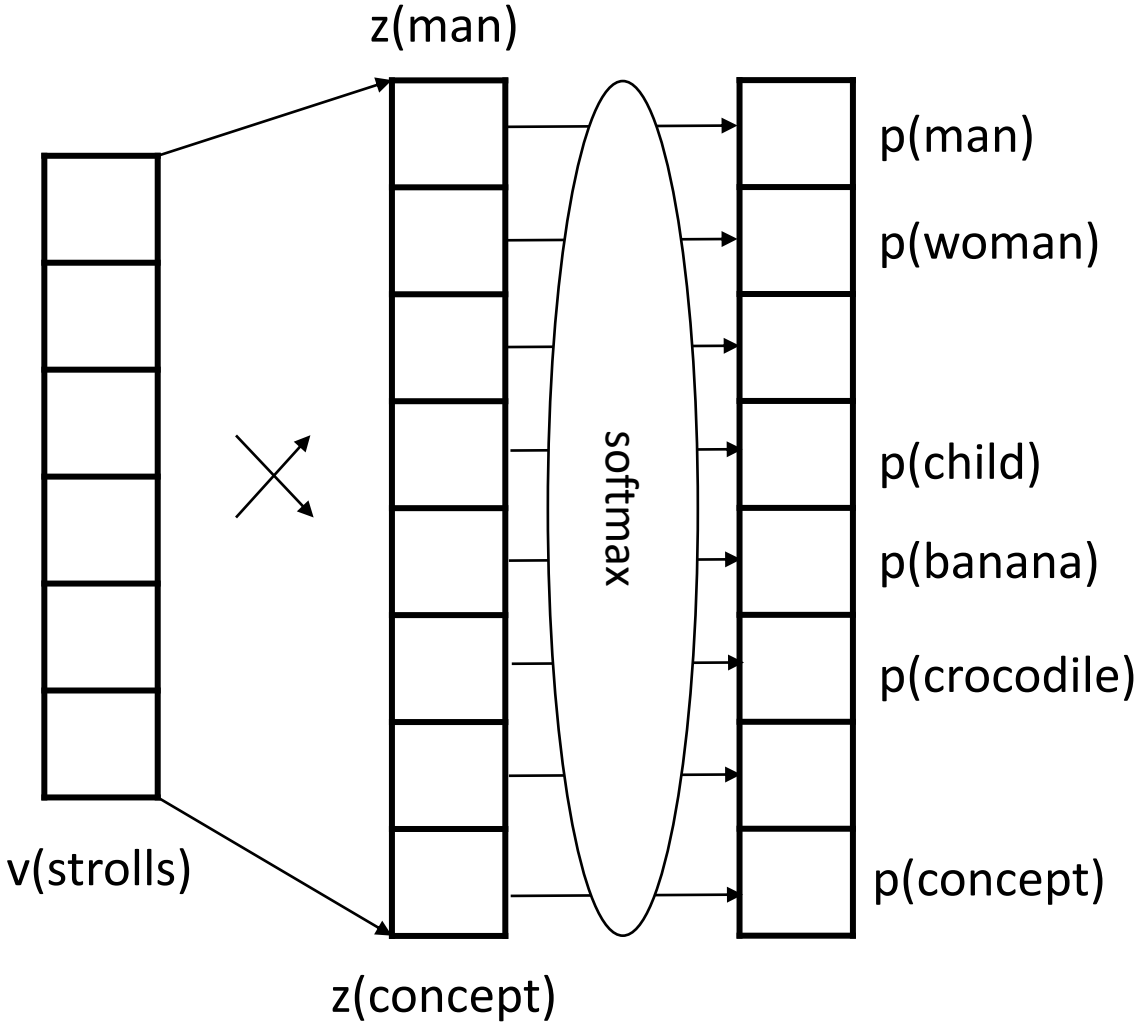

Predicting context words based on input words

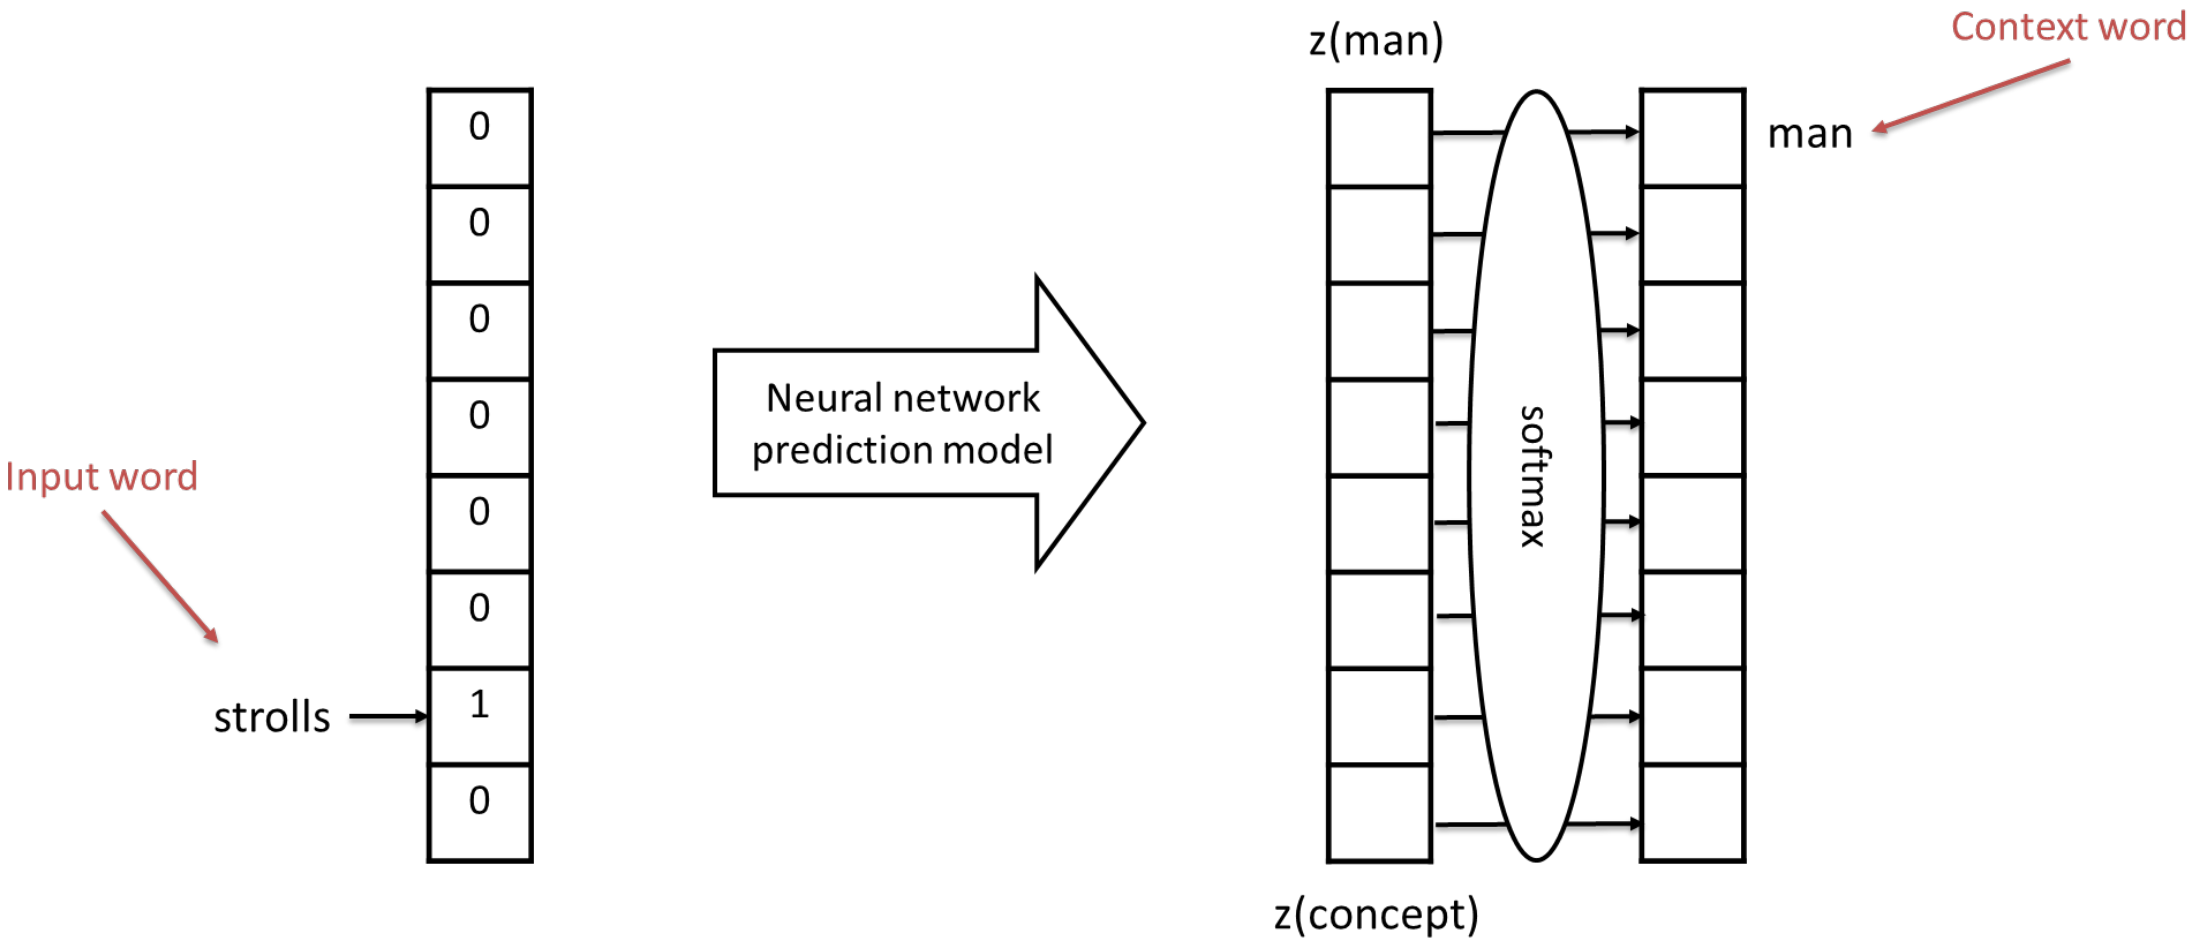

Input words and context words are one-hot encoded
(similar to bag of words representation)

Predicting context words based on input words

Training Data:

HUGE number of pairs of the following form:

{input word, context word}

e.g. from Wikipedia

Examples:

{strolls, man}
{strolls, woman}
{swims, crocodile}
{swims, fish}
{flies, bird}
{flies, plane}

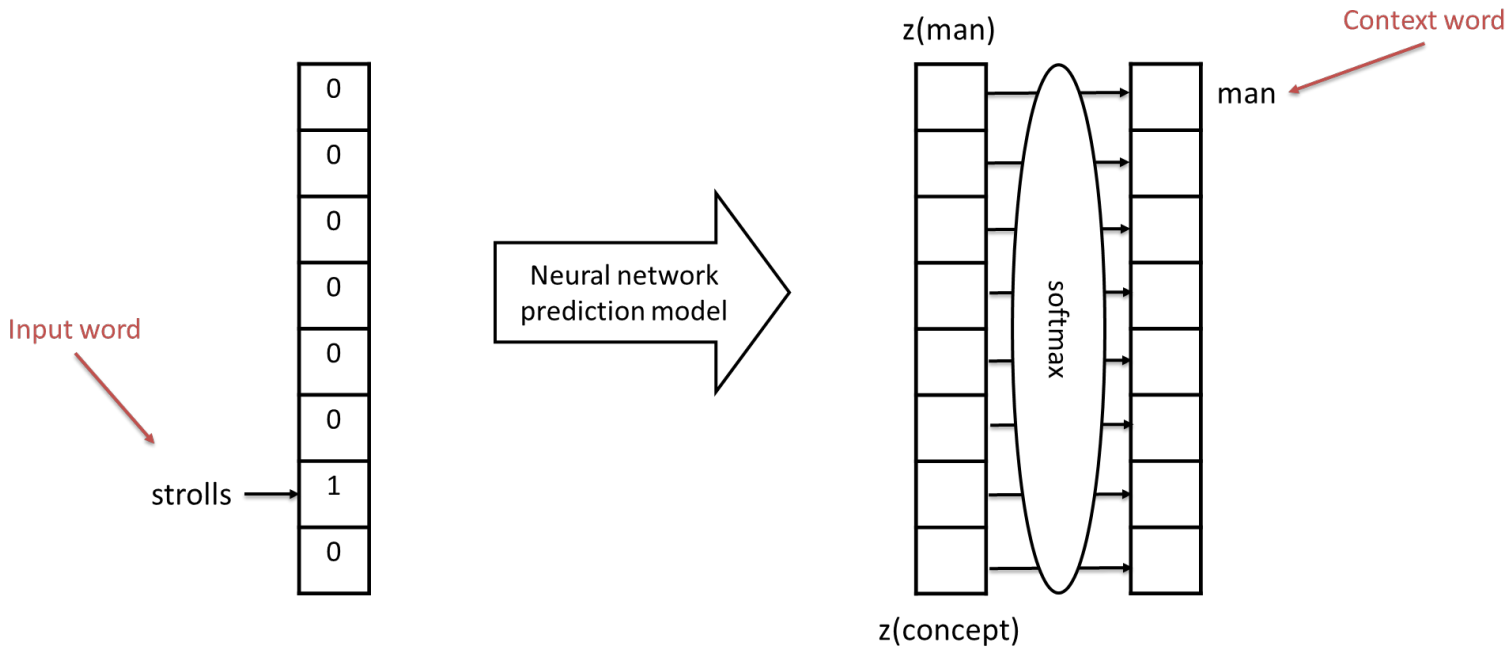

Predicting context words based on input words

Training Data:

HUGE number of pairs of the following form:

{input word, context word}

e.g. from Wikipedia

Examples:

{strolls, man}
{strolls, woman}
{swims, crocodile}
{swims, fish}
{flies, bird}
{flies, plane}

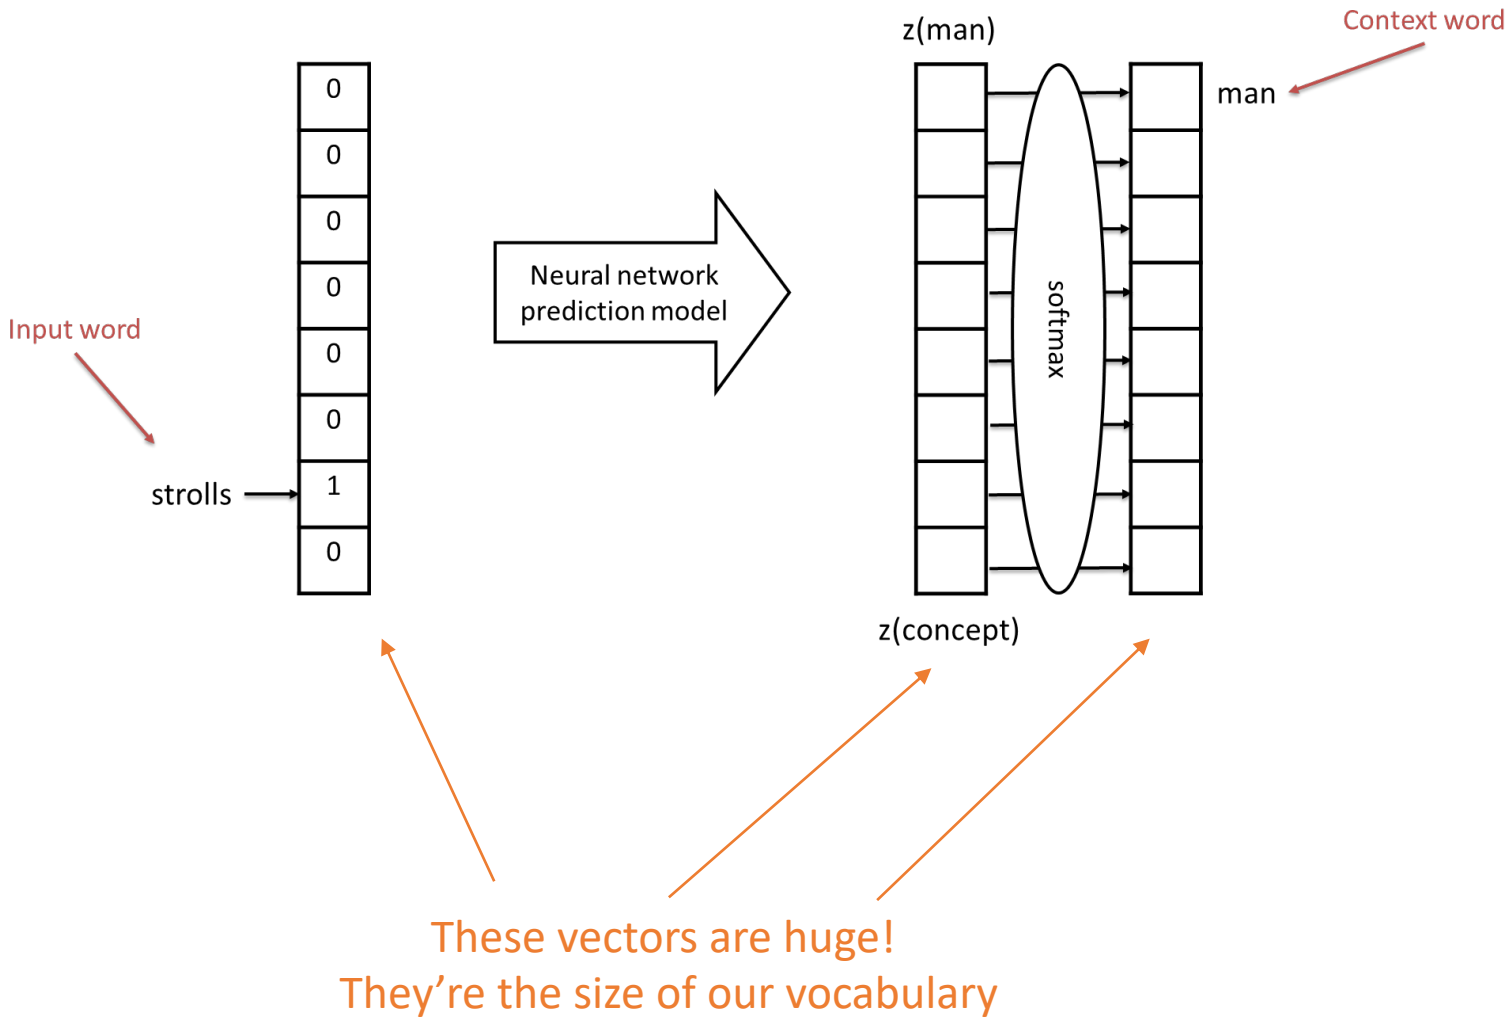

What's the simplest model we can possibly use?

First idea:

Directly connect our input to the log-odds layer

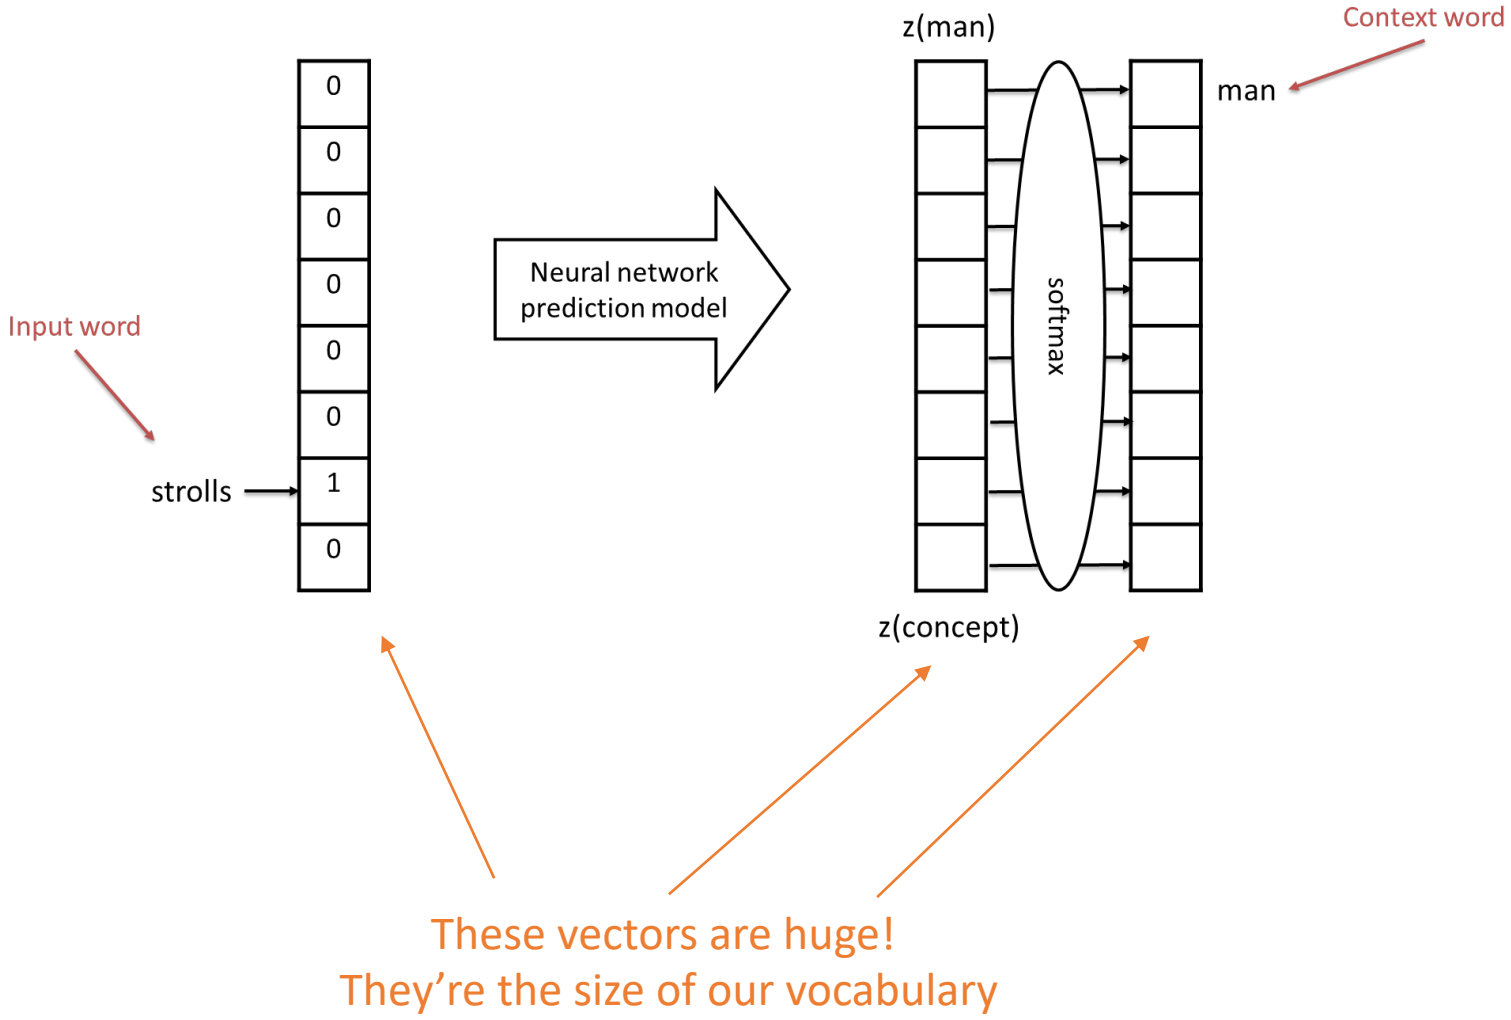

What's the simplest model we can possibly use?

First idea:

Direct connection from our input to the next-odd layer

How many connections?

V x V

Where V is our vocabulary size (approx. 6 billion)

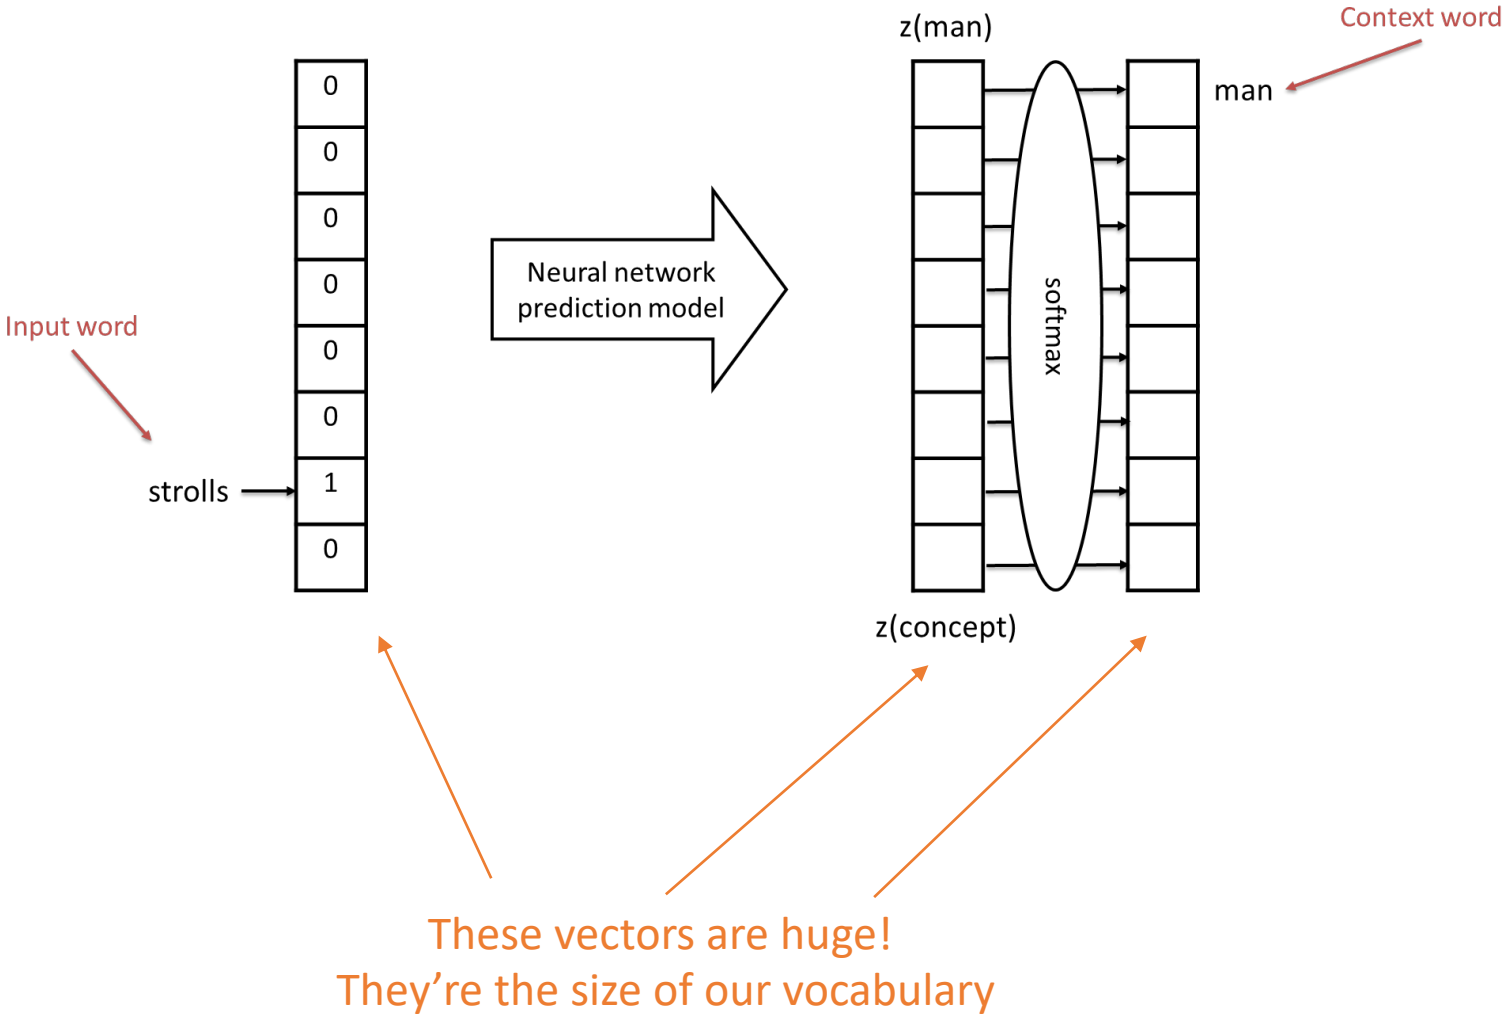

What's the next simplest?

How about a single hidden layer?

How many connections?

V x H x 2

Where V is our vocabulary size (approx. 6 billion)

And H is our hidden layer size ($\ll V$)

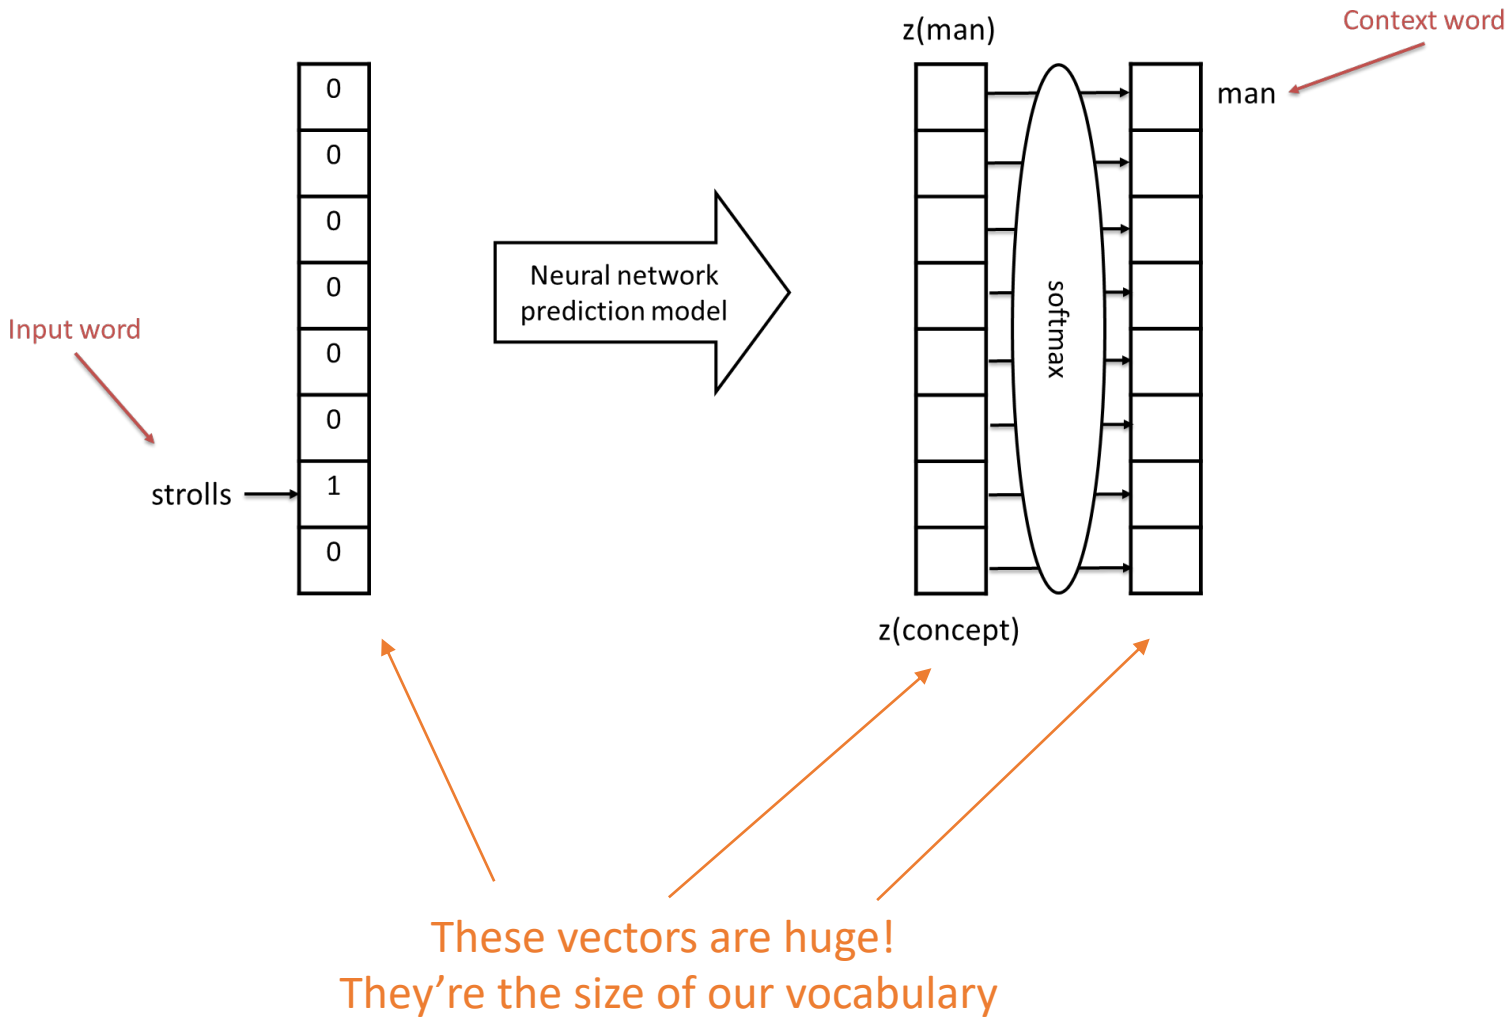

OK, let's try it: use a single hidden layer

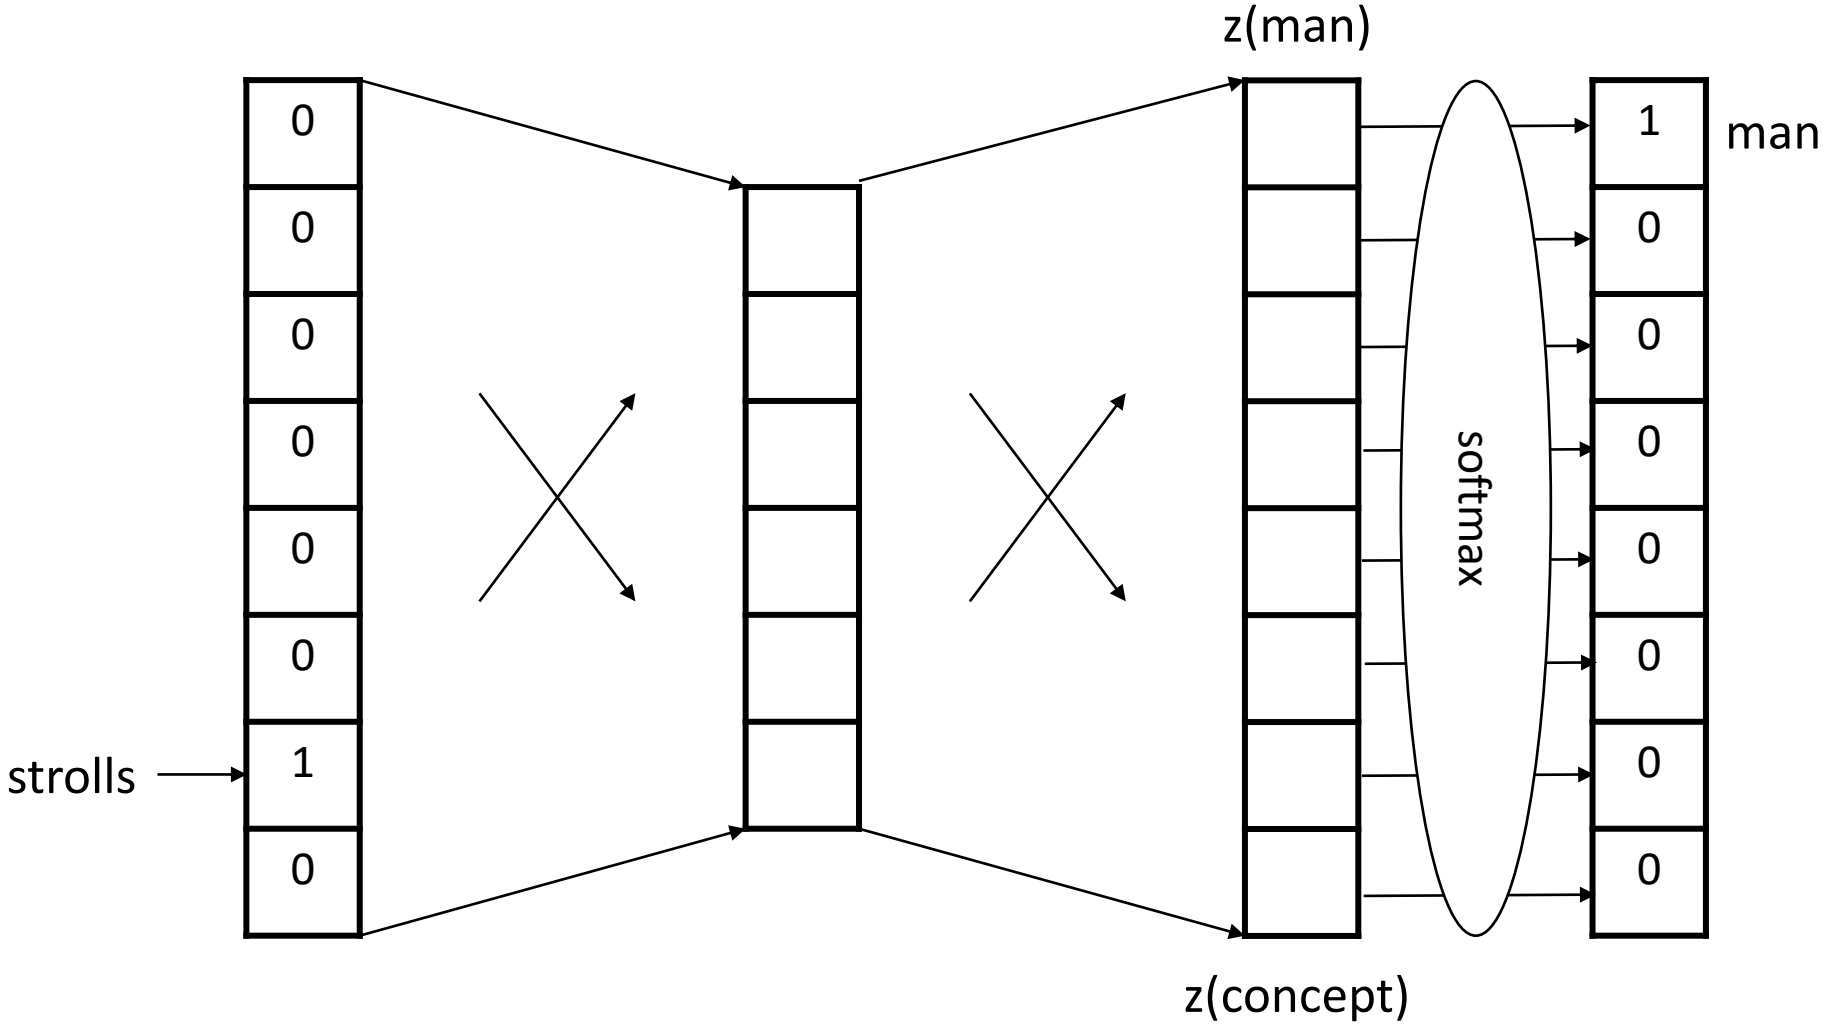

Use mini-batches of training examples; minimize cross-entropy loss

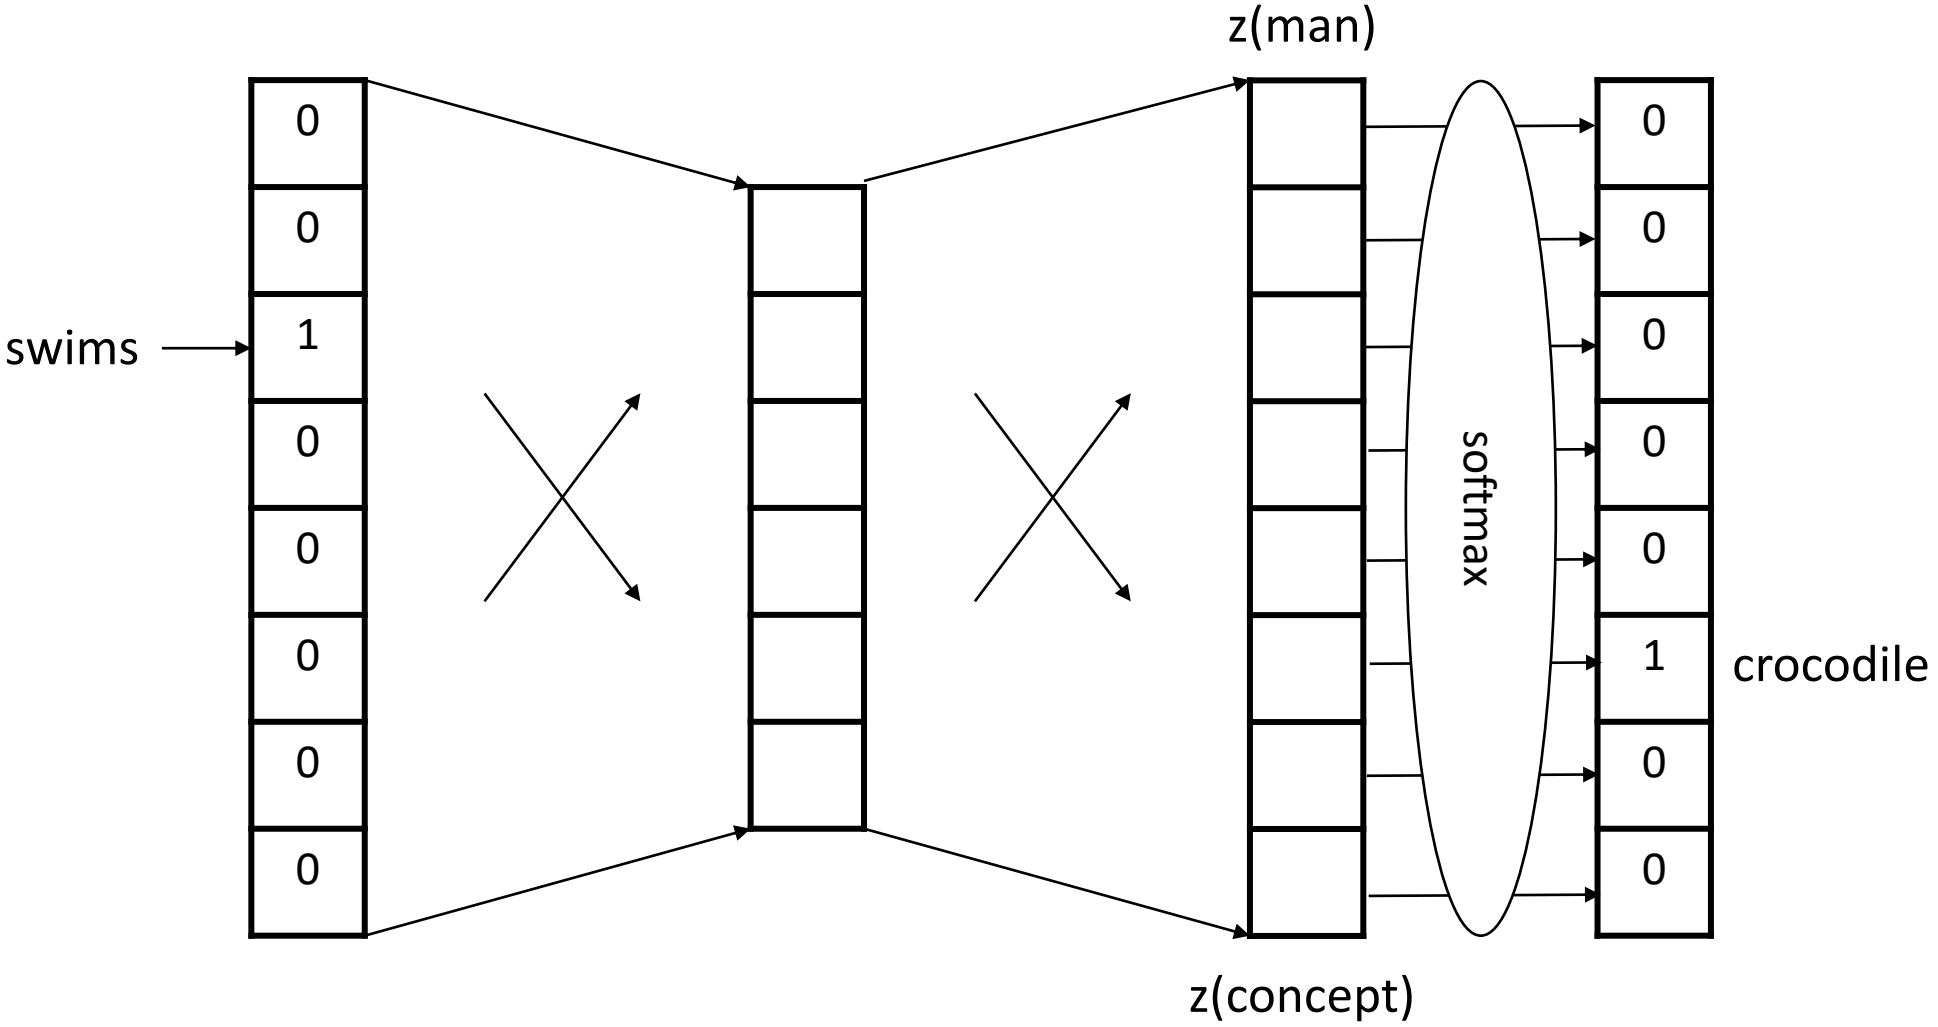

Learn our parameters: Weight Matrices W and B

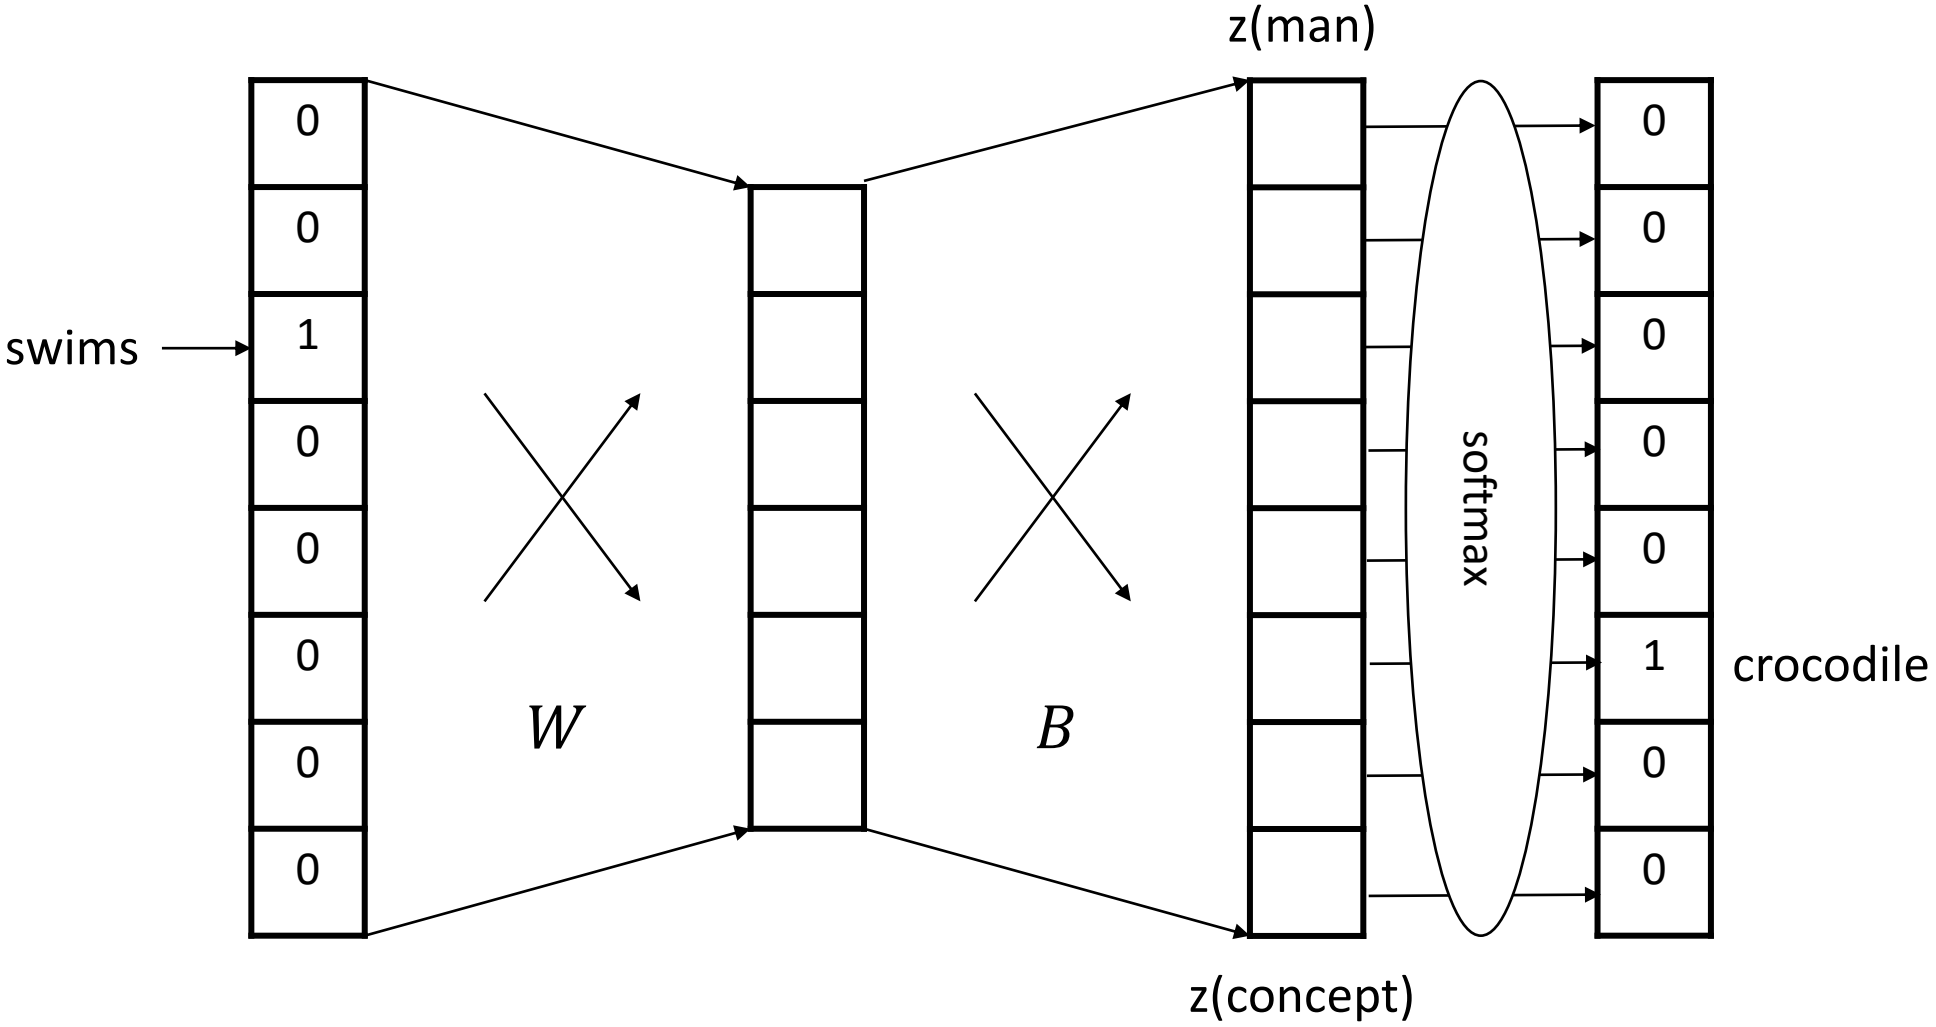

Isn't **this** the vector we were looking for?

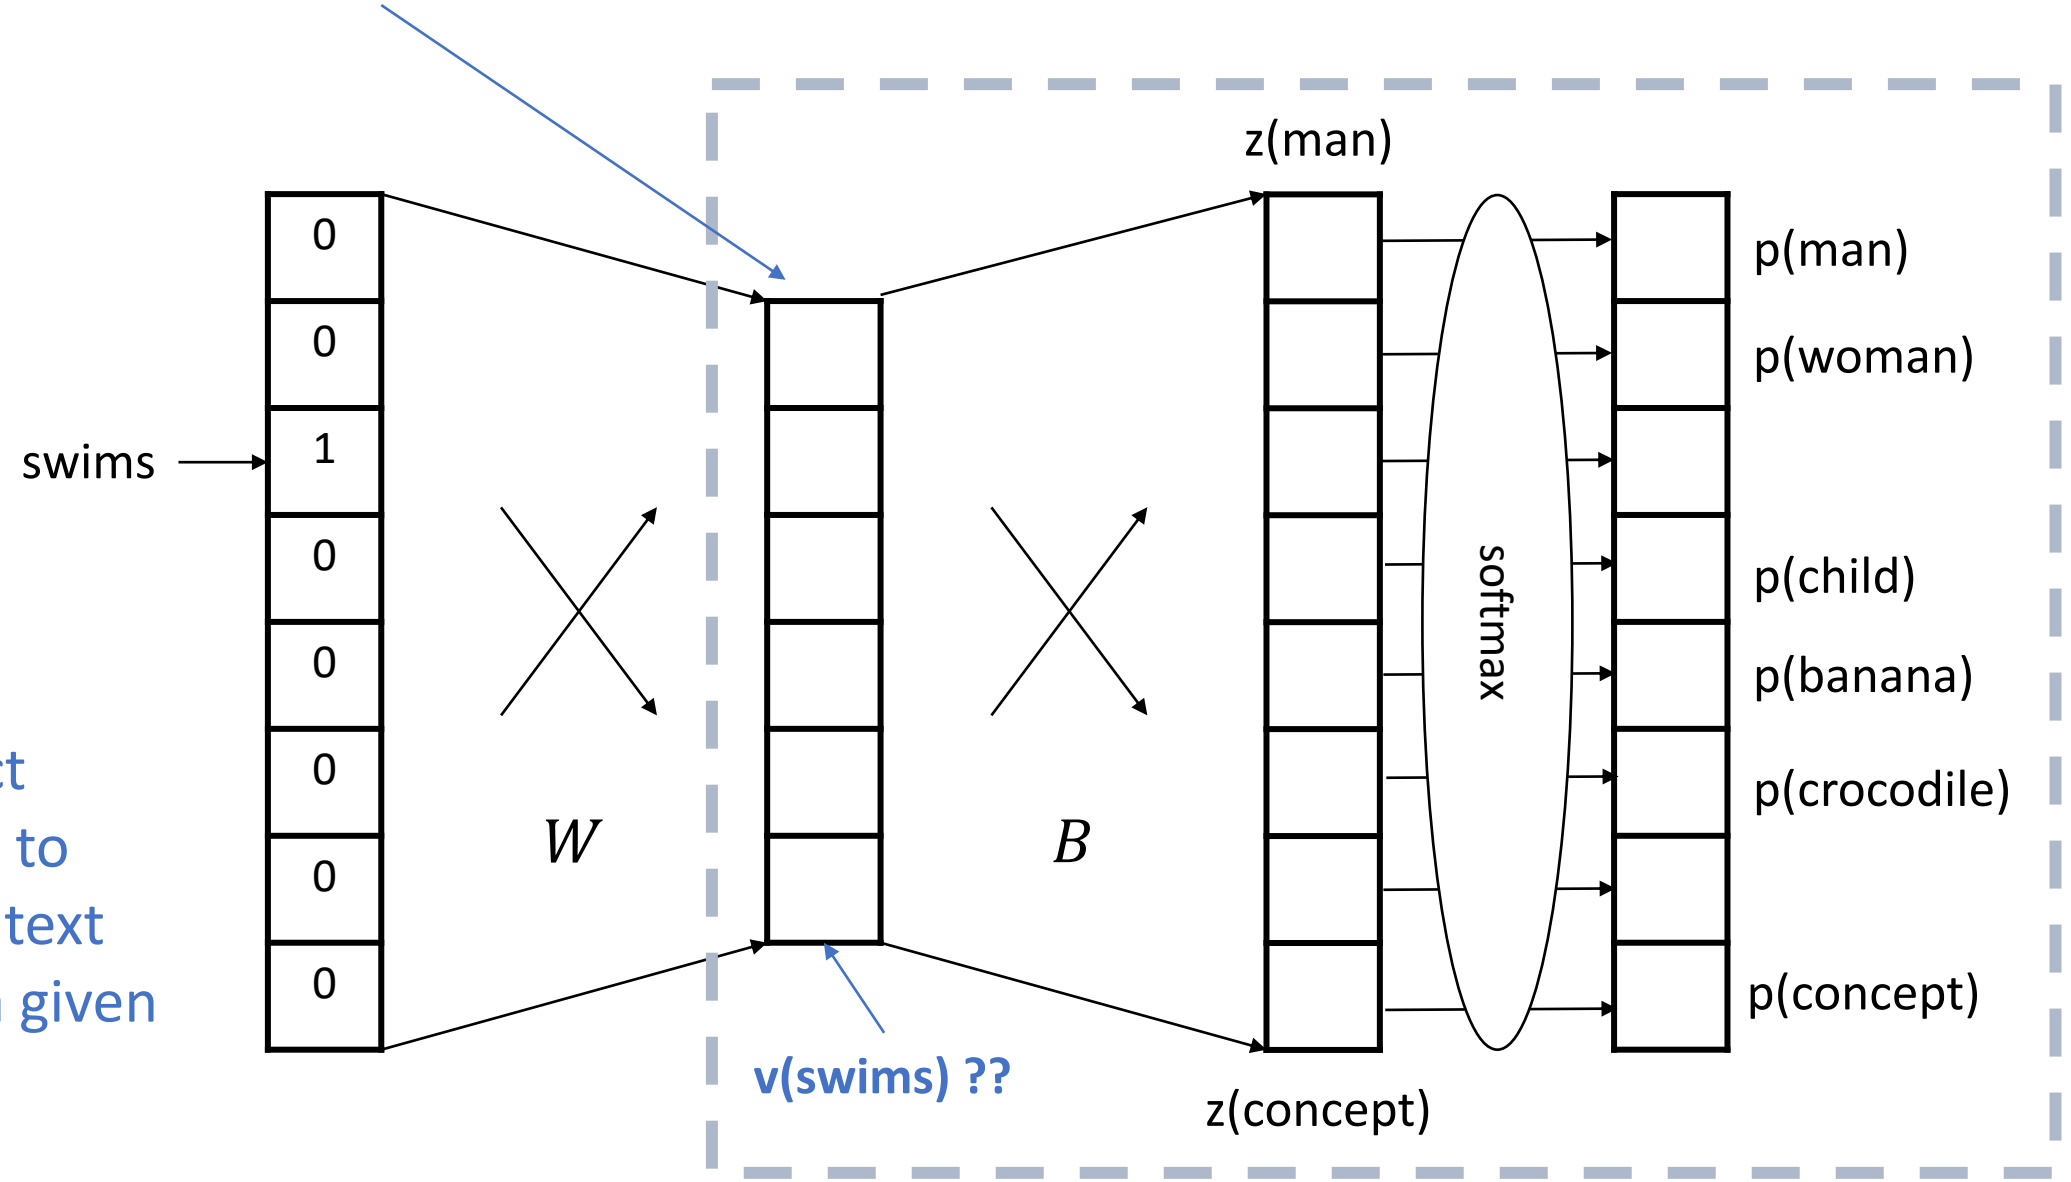

- it's compact
- It allows us to predict context words for a given input word

Let's take a closer look at W

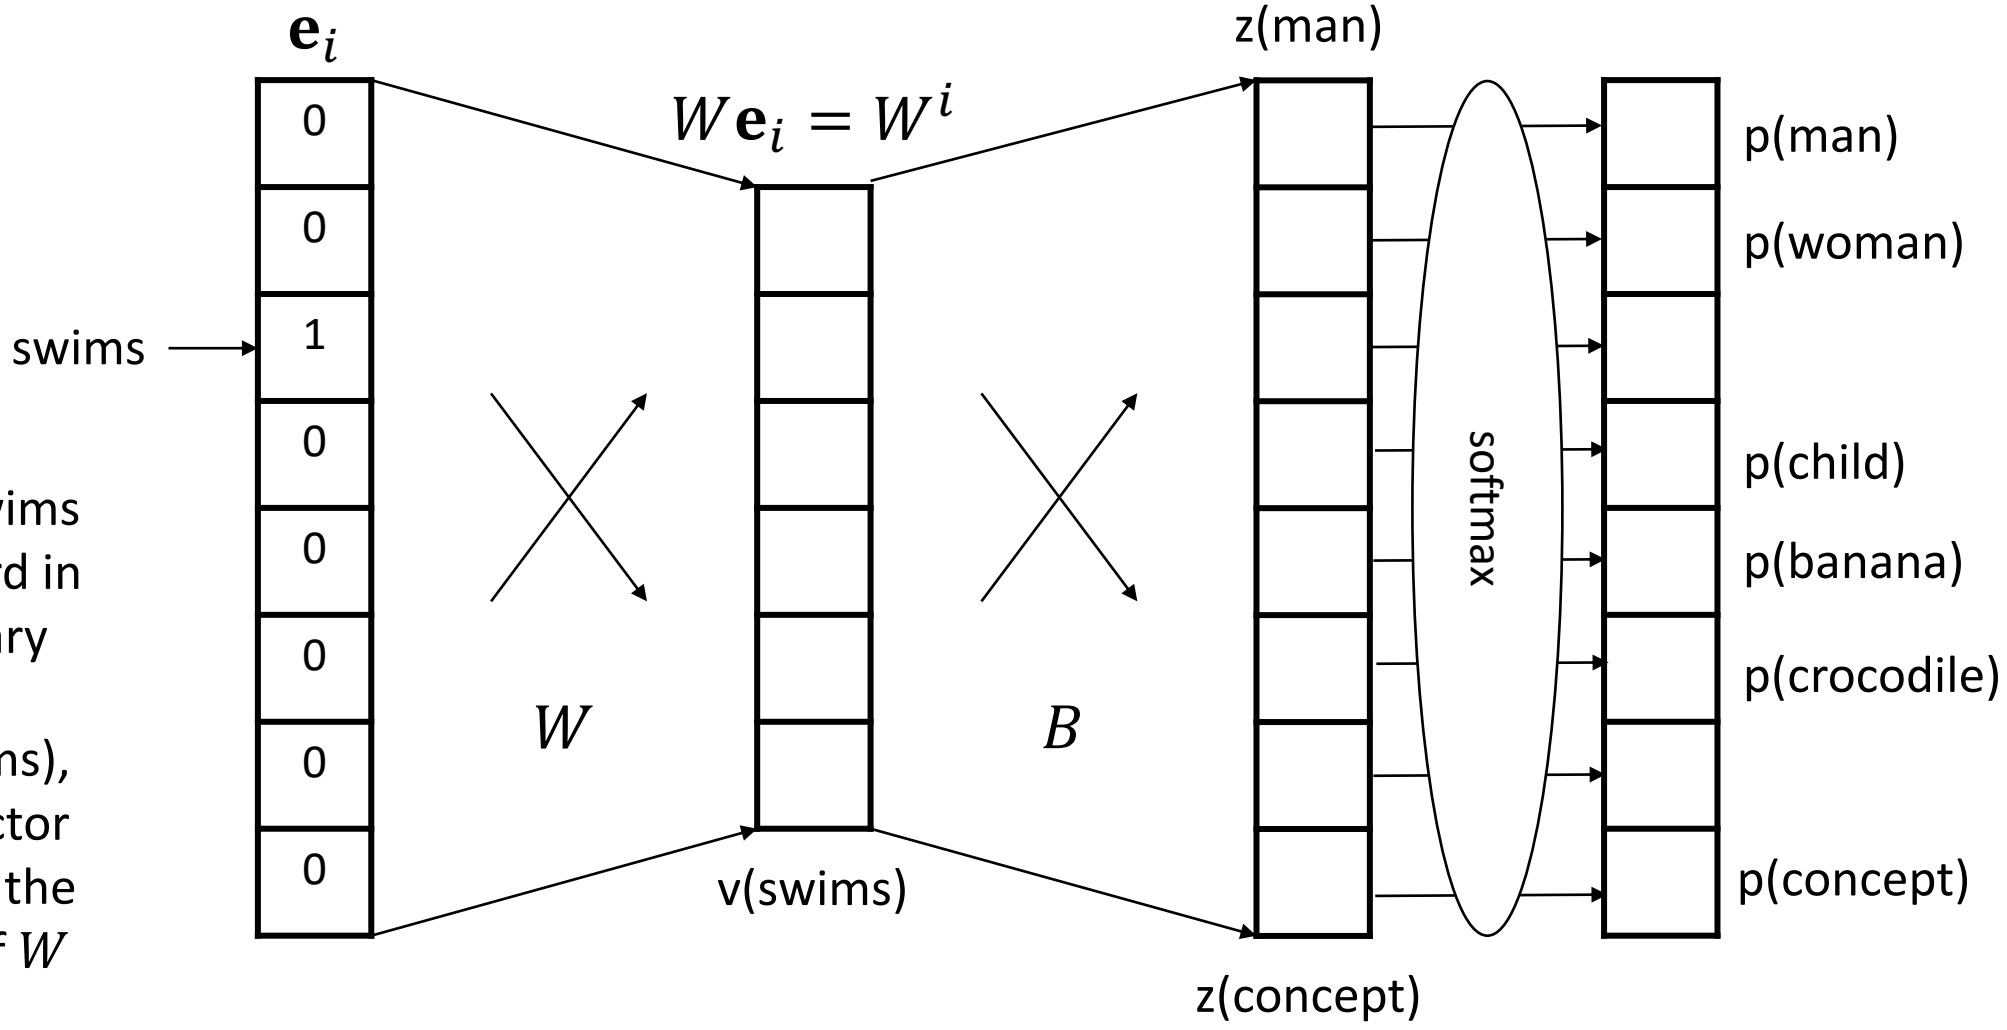

- suppose swims is the i^{th} word in our vocabulary

- then v(swims), the word vector for swims, is the i^{th} column of W

Let's take a closer look at W

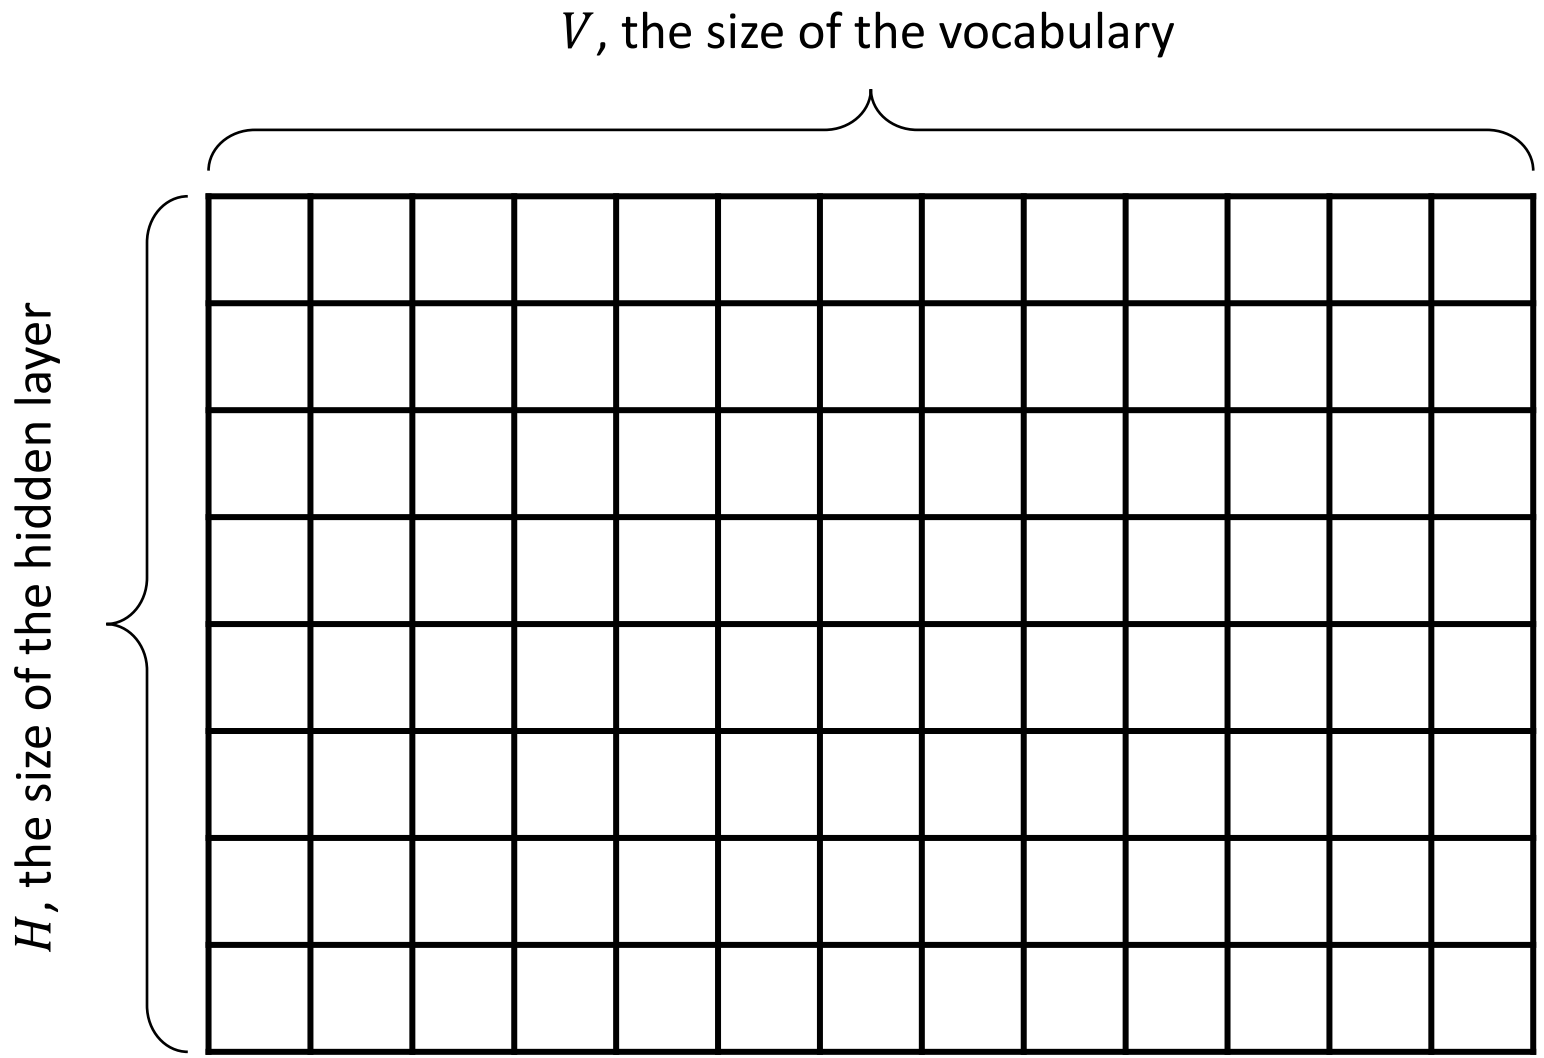

Let's take a closer look at W

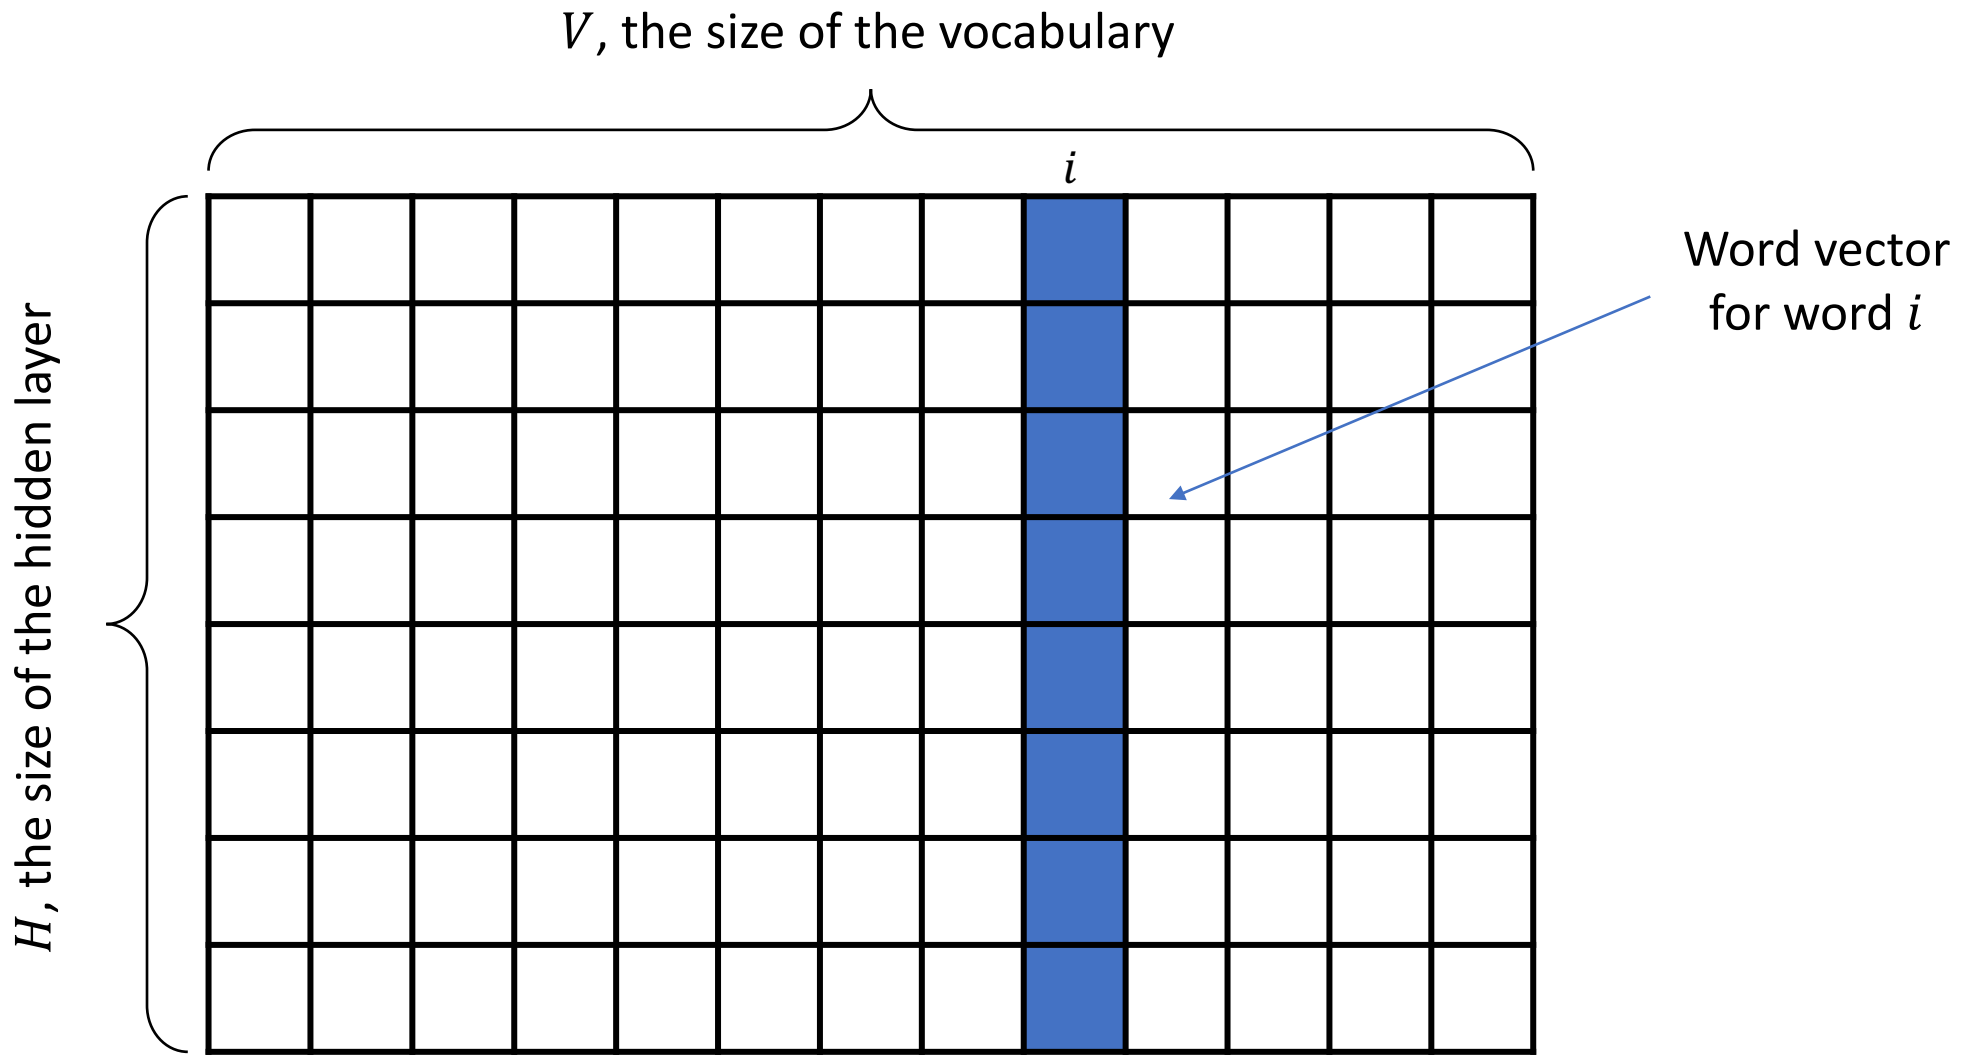

We now have a distributed representation of word *meaning* based on *context*

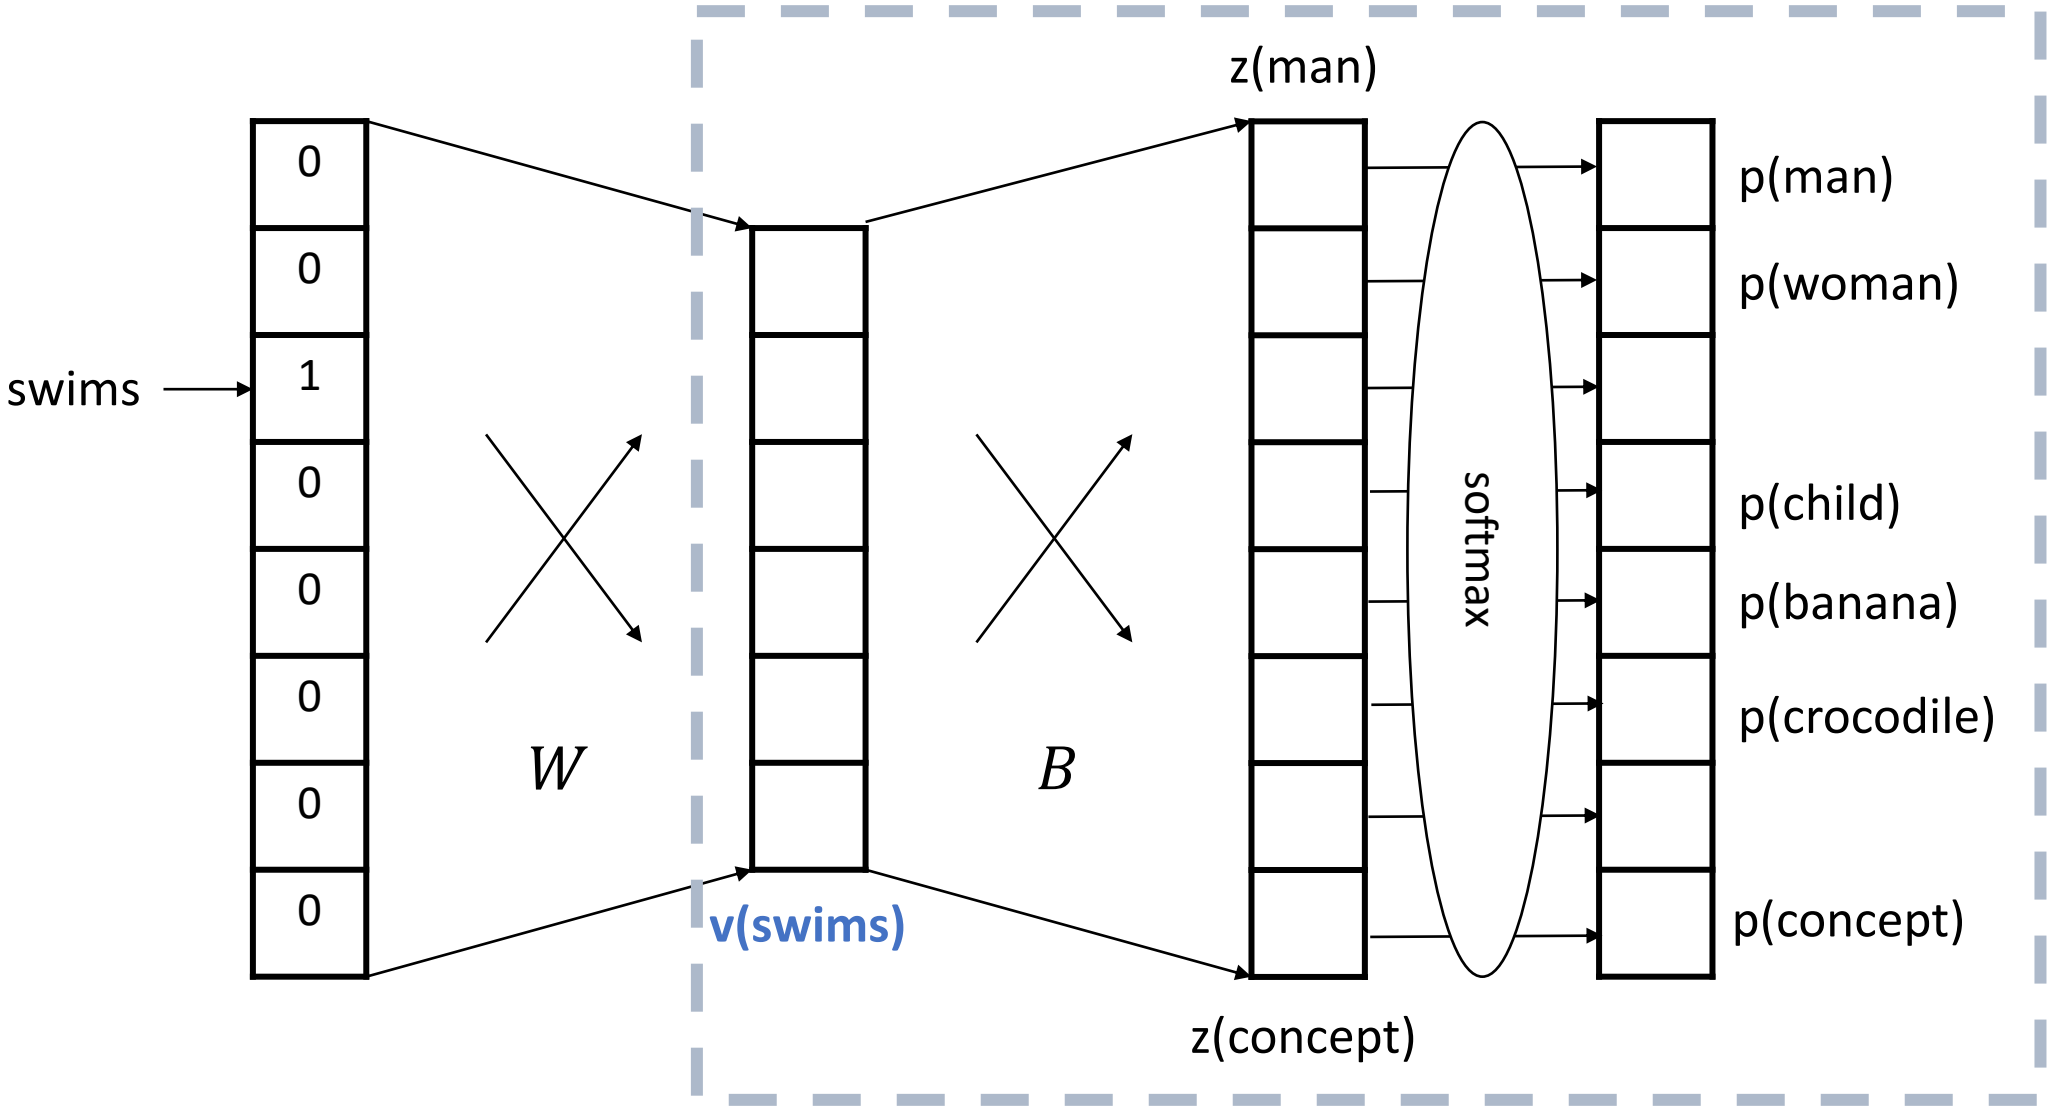

Important Takeaways:

- We are learning a vector representation for each word based on the contexts in which it appears
- training data: large number of pairs of nearby words from a large corpus
- These vectors give us much more flexibility when modeling: makes text sequences like other sequences