

# Model Learning, Validation, and Overfitting

June 11, 2020

Lecture 3, Applied Data Science  
MMCi Term 4, 2020

Matthew Engelhard

# MODEL LEARNING

Learning (or training) our model means:


-> Setting our parameters to the specific values that are the best match for our training data

What do we mean by “best match”?

-> We set our parameters to the values that maximize some measure of fit

-> Alternatively, we set them to values that minimize a penalty (i.e. a *loss*) that we choose

# Linear Regression Measure of Fit: Mean Squared Error





Error:

$$\epsilon_i = y_i - \hat{y}_i$$



Mean square error (MSE)

$$\frac{1}{N} \sum_{i=1}^N \epsilon_i^2$$

Suppose the slope is known.  
What happens to the MSE as we move the intercept ( $b_0$ )?



What is the best choice of intercept ( $b_0$ ) for these data, the one that minimizes the mean squared error?




# Logistic Regression Measure of Fit:

Probability that events  $y_1, \dots, y_N$  would have occurred if our model were correct


Probability that event  $y_i$  would have occurred if our model were correct:

$$(\sigma(z_i))^{y_i} (1 - \sigma(z_i))^{(1-y_i)}$$



# Logistic Regression Measure of Fit:

Probability that events  $y_1, \dots, y_N$  would have occurred if our model were correct



Probability that event  $y_i$  would have occurred if our model were correct:

$$(\sigma(z_i))^{y_i} (1 - \sigma(z_i))^{(1-y_i)}$$


A red bracket is drawn under the entire expression, with a red arrow pointing down to a piecewise definition:

$$\begin{cases} p_i, & y_i = 1 \\ 1, & y_i = 0 \end{cases}$$



# Logistic Regression Measure of Fit:

Probability that events  $y_1, \dots, y_N$  would have occurred if our model were correct



Probability that event  $y_i$  would have occurred if our model were correct:


$$(\sigma(z_i))^{y_i} (1 - \sigma(z_i))^{(1-y_i)}$$

A red bracket is drawn under the expression  $(\sigma(z_i))^{y_i} (1 - \sigma(z_i))^{(1-y_i)}$ . An arrow points from the bracket to a piecewise definition:

$$\begin{cases} 1, & y_i = 1 \\ 1 - p_i, & y_i = 0 \end{cases}$$

# Logistic Regression Measure of Fit:

Probability that events  $y_1, \dots, y_N$  would have occurred if our model were correct



Probability that event  $y_i$  would have occurred if our model were correct:

$$(\sigma(z_i))^{y_i} (1 - \sigma(z_i))^{(1-y_i)}$$

Probability that events  $y_1, \dots, y_N$  would have occurred if our model were correct:

$$\prod_{i=1}^N (\sigma(z_i))^{y_i} (1 - \sigma(z_i))^{(1-y_i)}$$

# Logistic Regression Measure of Fit:

Probability that events  $y_1, \dots, y_N$  would have occurred if our model were correct

Maximize the probability that events  $y_1, \dots, y_N$  would have occurred if our model were correct:

$$\prod_{i=1}^N (\sigma(z_i))^{y_i} (1 - \sigma(z_i))^{(1-y_i)}$$

Easier to maximize the log of this quantity (i.e. the log-likelihood):

$$\sum_{i=1}^N y_i \log \sigma(z_i) + (1 - y_i) \log(1 - \sigma(z_i))$$

# Maximize the log-likelihood = minimize the “cross-entropy” loss

Log-likelihood:


$$\sum_{i=1}^N y_i \log \sigma(z_i) + (1 - y_i) \log(1 - \sigma(z_i))$$

# Maximize the log-likelihood = minimize the “cross-entropy” loss

Cross-entropy loss:

$$\sum_{i=1}^N -y_i \log \sigma(z_i) - (1 - y_i) \log(1 - \sigma(z_i))$$

The cross-entropy loss is also used for other classification models, including convolutional neural network classifiers for image processing.



# *How* do we minimize the loss?

- The cross-entropy loss just tells us *what* quantity we should be minimizing
- In some cases (e.g. linear regression), we can solve for the minimum directly
- But, we'd like to have an approach that works even for very complex models

# Strategy: determine how small changes in parameters affect the loss

MORTALITY PREDICTION WORKSHEET												
COVARIATES							OUTCOMES AND PREDICTIONS					
patient	age	age_normalized	female	temp	temp_normalized		mortality	predicted_log_odds	predicted_prob	prediction	correct?	loss
0	30.5	-0.5	0	105.0	2.4		1	0.00	0.50	0	0	0.3010
1	74.0	1.1	1	96.7	-0.8		0	0.00	0.50	0	1	0.3010
2	27.4	-0.6	0	96.1	-1.0		0	0.00	0.50	0	1	0.3010
3	0.1	-1.5	1	98.5	-0.1		0	0.00	0.50	0	1	0.3010
4	0.7	-1.5	1	96.5	-0.9		0	0.00	0.50	0	1	0.3010
5	49.9	0.2	1	97.1	-0.6		0	0.00	0.50	0	1	0.3010
6	72.9	1.0	1	100.1	0.5		1	0.00	0.50	0	0	0.3010
7	29.1	-0.5	1	99.6	0.3		0	0.00	0.50	0	1	0.3010
8	83.5	1.4	1	100.6	0.7		1	0.00	0.50	0	0	0.3010
9	82.3	1.4	1	95.2	-1.3		1	0.00	0.50	0	0	0.3010
10	23.7	-0.7	0	99.4	0.2		1	0.00	0.50	0	0	0.3010
11	12.9	-1.1	0	96.6	-0.8		0	0.00	0.50	0	1	0.3010
12	53.9	0.4	1	100.3	0.6		0	0.00	0.50	0	1	0.3010
13	18.8	-0.9	0	98.6	0.0		0	0.00	0.50	0	1	0.3010
14	51.8	0.3	0	98.5	-0.1		0	0.00	0.50	0	1	0.3010
15	3.3	-1.4	0	94.6	-1.6		0	0.00	0.50	0	1	0.3010
16	69.7	0.9	0	99.1	0.1		0	0.00	0.50	0	1	0.3010
17	60.4	0.6	1	104.2	2.1		1	0.00	0.50	0	0	0.3010
18	73.6	1.1	1	99.1	0.1		1	0.00	0.50	0	0	0.3010
19	53.3	0.3	1	99.1	0.1		0	0.00	0.50	0	1	0.3010
PARAMETERS		b_age	b_female		b_temp	bias				PERFORMANCE	accuracy	avg_loss
	guess	0.00	0.00		0.00	0.00					0.65	0.3010
	optimal											

# Strategy: determine how small changes in parameters affect the loss


We know:

- If we rotate G1 by 1 radian, G2 will rotate by  $-12/8$  radians.
- If G2 rotates by 1 radian, G3 will rotate by  $-8/6$  radians.

How do we determine the effect of G1 on G3?

➤ Multiply the effects.

➤  $\left(-\frac{12}{8}\right) * \left(-\frac{8}{6}\right) = \frac{12}{6} = 2$





# Strategy: determine how small changes in parameters affect the loss


We know:

- If we rotate G1 by 1 radian, G2 will rotate by  $-12/8$  radians.
- If G2 rotates by 1 radian, G3 will rotate by  $-8/6$  radians.

How do we determine the effect of G1 on G3?

➤ Multiply the effects.

➤  $\left(-\frac{12}{8}\right) * \left(-\frac{8}{6}\right) = \frac{12}{6} = 2$




# Strategy: determine how small changes in parameters affect the loss

We know:

- If we increase  $b_1$  by a small amount  $\varepsilon$ , then  $z_i$  will increase by  $\varepsilon * x_{i1}$
- If we increase  $z_i$  by a small amount  $\varepsilon$ , then  $p_i$  will increase by  $\varepsilon * \frac{d\sigma(z_i)}{dz_i}$  (depends on  $z_i$ )

How do we determine the effect  $b_1$  on the cross-entropy loss (which depends on  $p_i$ )?


➤ Multiply the effects.



# Strategy: determine how small changes in parameters affect the loss

This is called the *chain rule* (calc 101)


- We use it to see how small changes in parameters affect the loss
- Could be a very long chain...
- Some parameters have a greater effect than others
- We change all parameters at once, with each change proportional to that parameter's effect on the loss
  - This is *gradient descent*



# Strategy: determine how small changes in parameters affect the loss

It could be a very long (and complex) chain...

- If we increase  $x_{i1}$  by  $\varepsilon$ , it changes *all* of the  $\eta_{ij}$ ...
- ...each of which changes *all* of the  $\zeta_{ij}$
- ...each of which changes  $p_i$
- Machine learning software like TensorFlow allows us to keep track, even for very complicated models





# *Learning*: find parameters that minimize the loss

Let's begin by finding the value of a *single* parameter that minimizes the loss. We'll consider the intercept  $b_0$  of a linear regression model.

$$y_i = bx_i + b_0 + \varepsilon_i$$

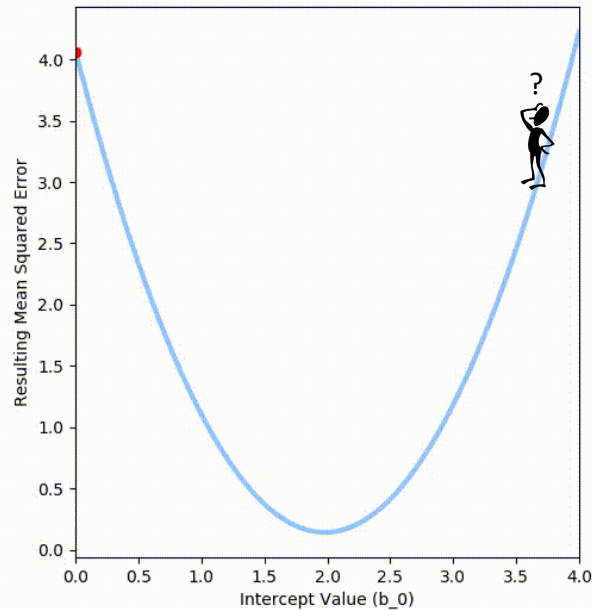

# *Learning*: find parameters that minimize the loss

Easy!!



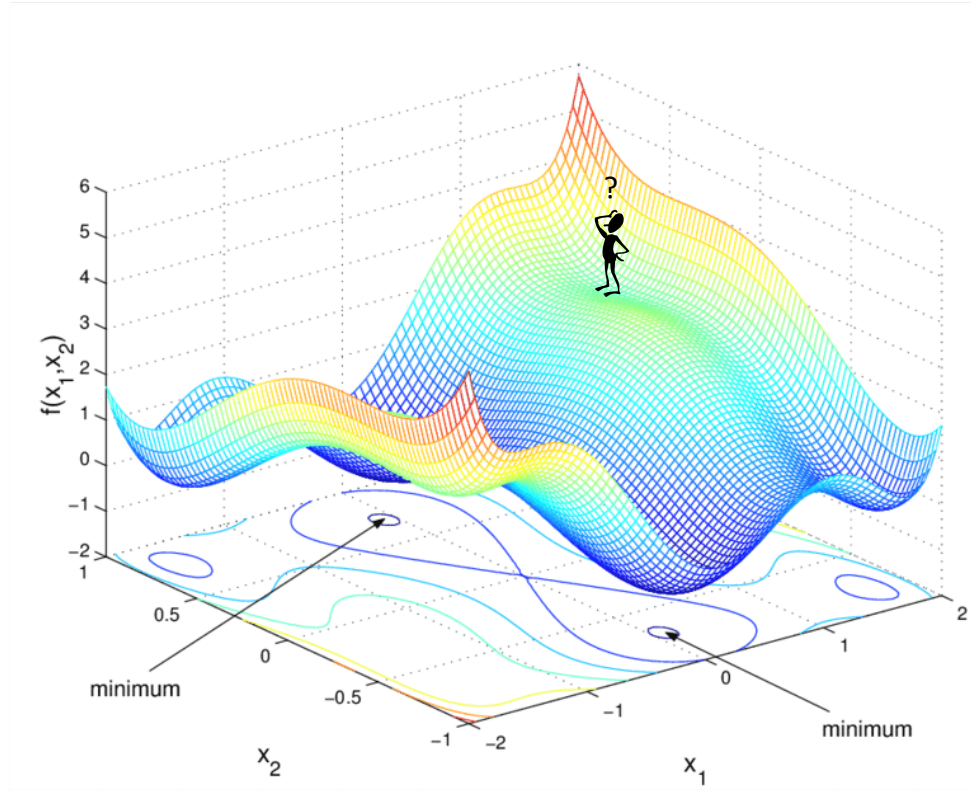
*Learning:* find parameters that minimize the loss

Easy!!



*Learning*: find parameters that minimize the loss

Easy?





*Learning:* find parameters that minimize the loss

Not easy.



*Learning:* find parameters that minimize the loss

Not easy.




# Learning: find parameters that minimize the loss

- With deep learning models, we are trying to minimize a function of many variables
- Can't visualize it
- Can't solve for the minimum directly
- So, we follow the slope and hope for the best (i.e. gradient descent)
- May end up in a low point that isn't the lowest, i.e. *local minimum*
- But, if we have lots of data, things usually work out OK




# OVERFITTING

# We like flexible, non-linear decision boundaries...



# But some models can be *too* flexible.



Green boundary:

- This is overfitting

Black boundary:


- Balance between fit and model complexity

-> The black boundary is likely to perform better on new data




# Overfitting

“Overfitting” happens when the learned model increases complexity to fit the observed training data *too well* – will not work to predict future data!

What would we want to use to fit these example data points?






# Classic Example: Increasing Polynomial Order



Increasing complexity does not seem appropriate...

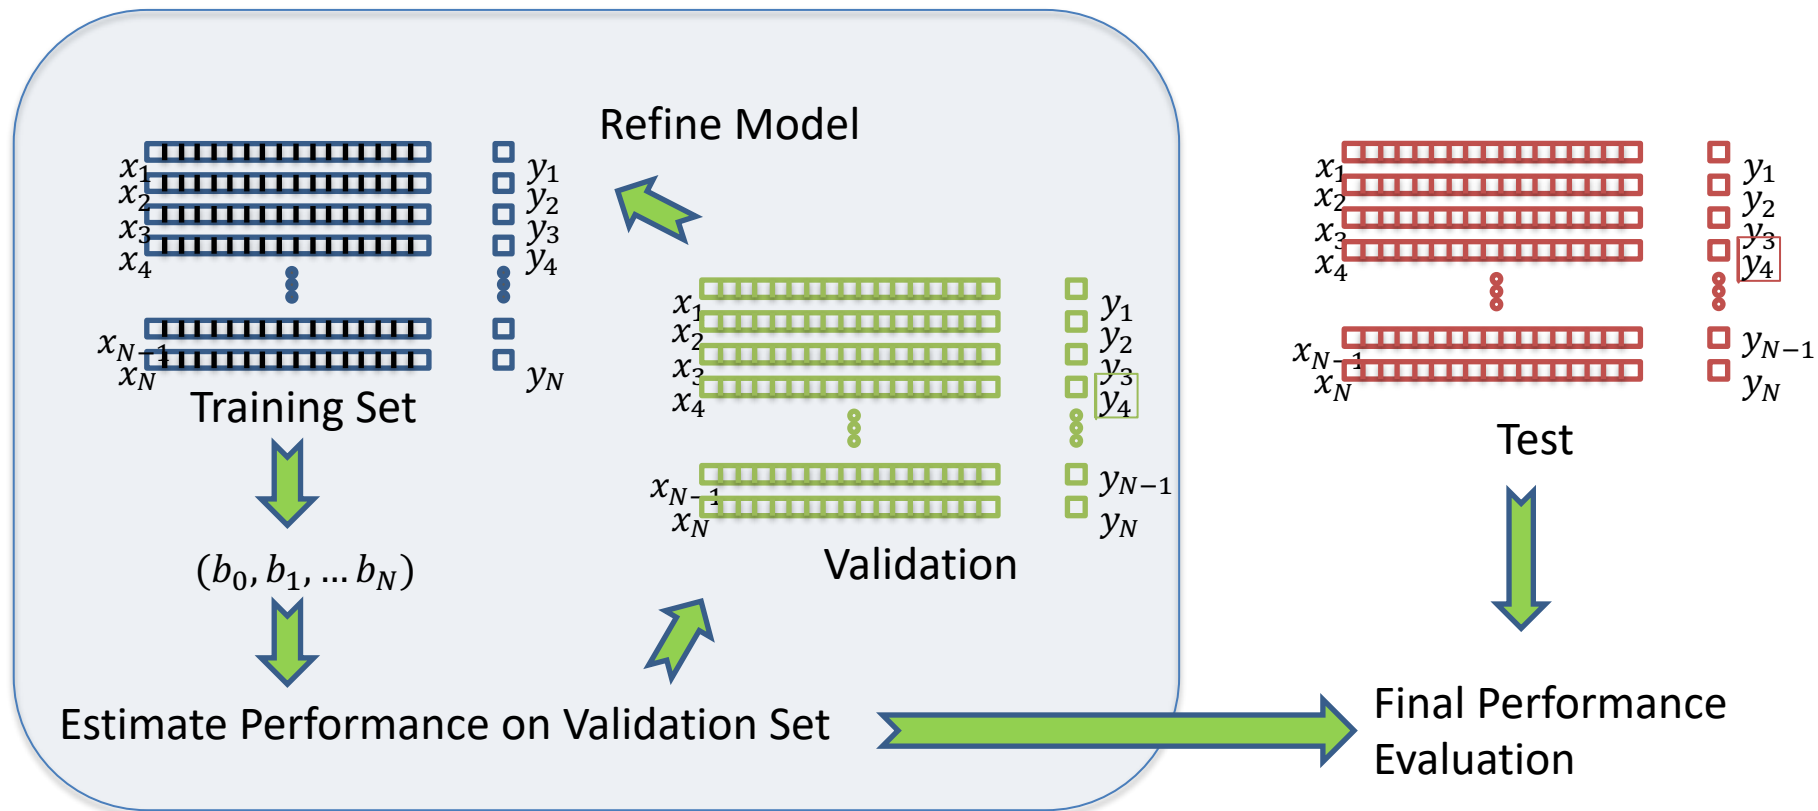


# With a flexible enough model, we can typically reach 100% accuracy on our training set



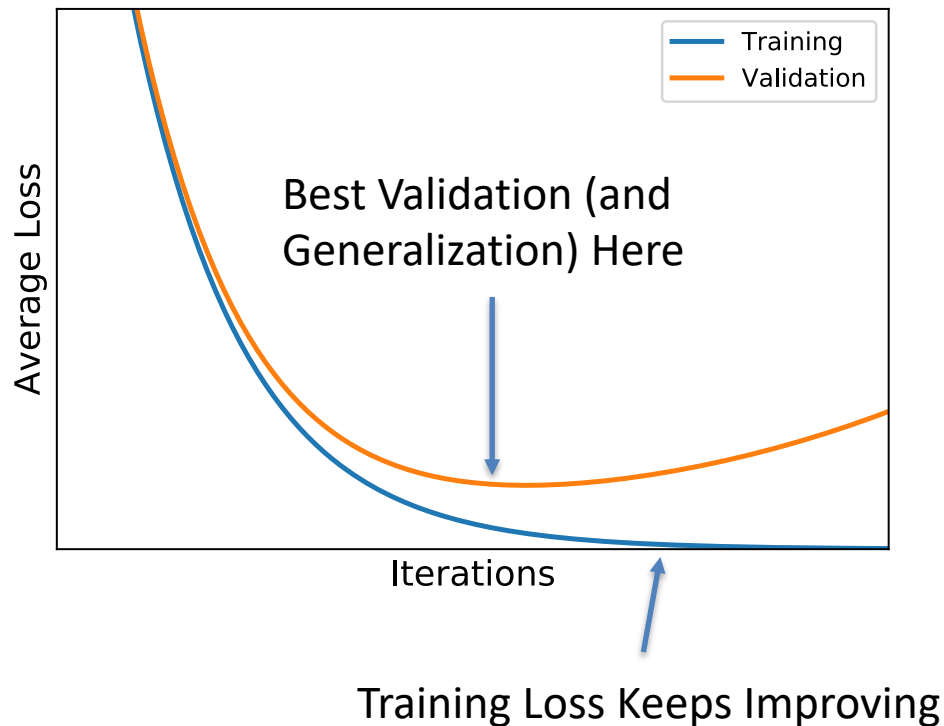
Increasing complexity does not seem appropriate...

# Using a Test Set Protects Against Overfitting: See how the model performs on *new* data



## Early Stopping

- During optimization, we can check the validation loss as we go.
- Instead of optimizing to convergence, we can optimize until the *validation* loss stops improving
  - Saves computational cost
  - Performs better on validation (and test) sets
- Widely used technique in the field



# Other Ways to Use the Validation Set

- Choosing between models (e.g. MLP, LR)
- Choosing how strongly to *regularize* the model, i.e. penalize parameters
- Choosing network depth, width, etc.
- Many other “hyperparameters” we might tune

# Conclusions

- *Learning* consists in setting model parameters to maximize some measure of fit (or equivalently, minimize some loss)
- This sounds easy, but is difficult in practice when working with complex models
- Greater model complexity is often, but not always, advantageous
- Proper model validation is critical to estimate real-world performance and prevent overfitting