

# Convolutional Neural Networks

June 7, 2019


Block 2, Lecture 1  
Applied Data Science  
MMCi Term 4, 2019

Matthew Engelhard

Many slides created by Tim Dunn

# Deep Learning for Image Analysis

## Diabetic Retinopathy Classification




Healthy Retina



Unhealthy Retina

# Deep Learning for Diabetic Retinopathy Classification




$$\text{sensitivity} = \frac{\text{number of true positives}}{\text{total number of positives in the dataset}}$$

$$\text{specificity} = \frac{\text{number of true negatives}}{\text{total number of negatives in the dataset}}$$

Gulshan et al. *JAMA* (2016)

# Deep Learning for Image Analysis


TSA  
Screening



Liang et al. *SPIE* (2018)

# Deep Learning for Image Analysis

## Markerless Motion Capture: Automatic 3D Surface Meshes from Video



DensePose  
(Facebook)

# Mask R-CNN



# Deep Learning for Image Analysis

## Style Transfer and Harmonization



Gatys et al. A Neural Algorithm of Artistic Style. *arXiv* (2015)



# Deep Learning for Image Analysis

## Style Transfer and Harmonization






Luan et al. Deep Painterly Harmonization. *arXiv* (2018)






# Motivating the CNN: Back to MNIST


Zeros



Ones

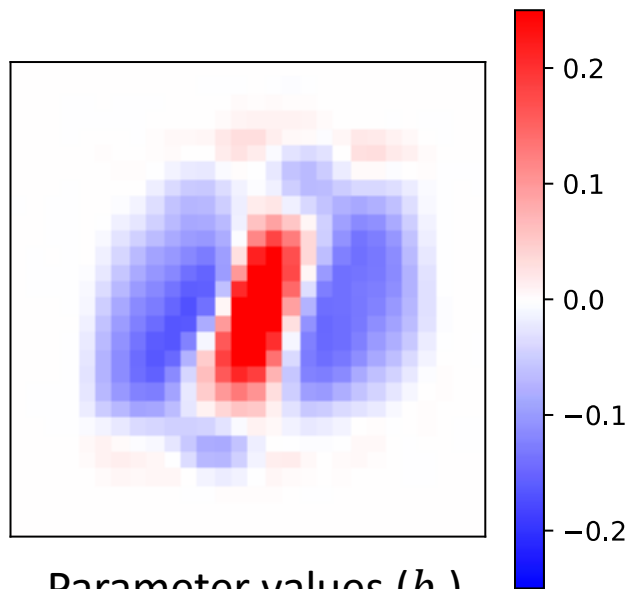


# A “Filter” to Detect Ones




Parameter values  $(b_j)$   
for each pixel  $j$

# A “Filter” to Detect Ones

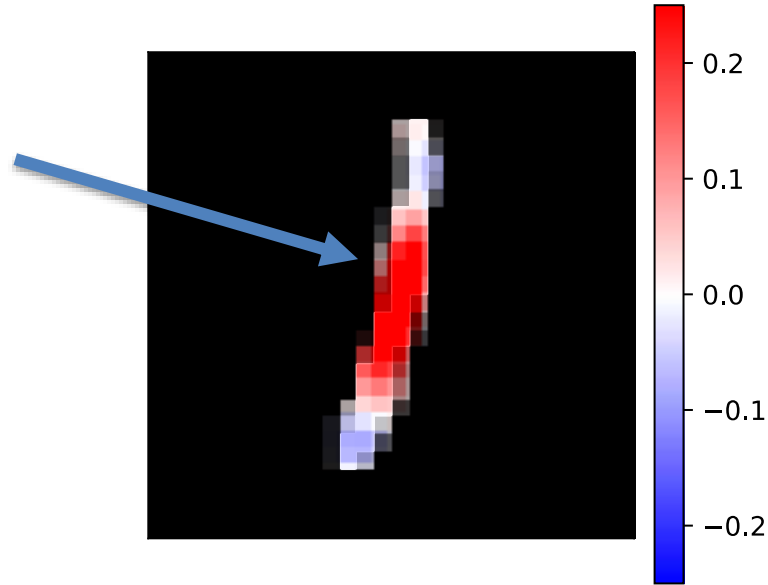


Parameter values ( $b_j$ )  
for each pixel  $j$




# A “Filter” to Detect Ones

Positive Section



# A “Filter” to Detect Ones




What if we'd like to find a 1 anywhere in a larger image?







# Searching for a 1...




# Searching for a 1...





# Searching for a 1...





Our previous approach looks for a 1 at a specific location.



Our previous approach looks for a 1 at a specific location.





If we move the position of the 1, it no longer works.






Instead of using logistic regression or an MLP, let's look for a new kind of model, one with more flexible filters




**"1"**


Instead of a filter that's the size of the whole image,  
we'd like a smaller filter that we can move around




As we move this filter, we calculate the inner product between the filter itself and the portion of the image that's underneath it.




When the filter is placed over a region that looks like the filter, the inner product (i.e. filter output) will be large.



When it's placed over a region that does not look like the filter, the inner product (i.e. filter output) will be small

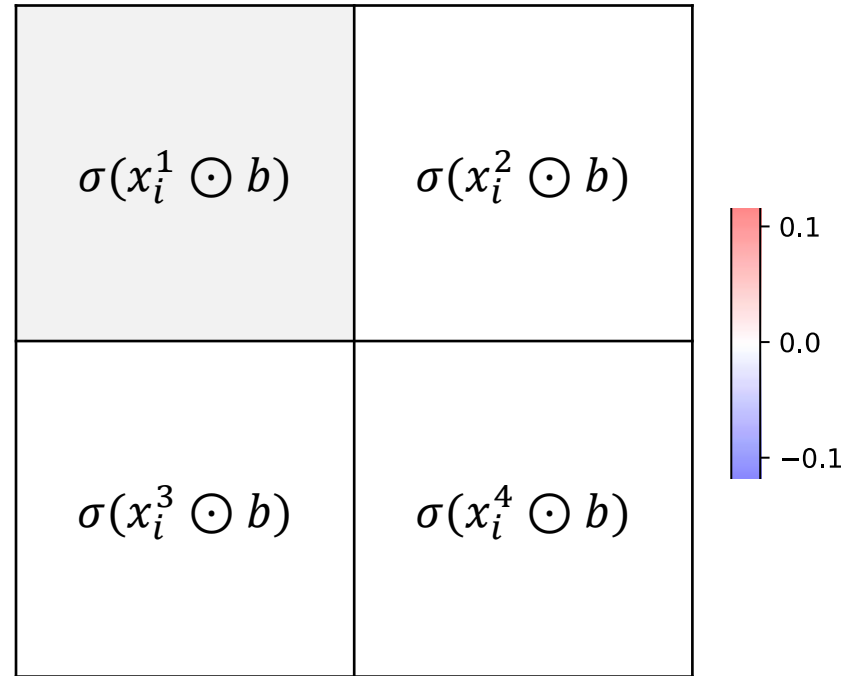



Examining filter output  $\sigma(x_i^R \odot b)$ , where  $x_i^R$  is the portion of image  $i$  where the filter is placed.

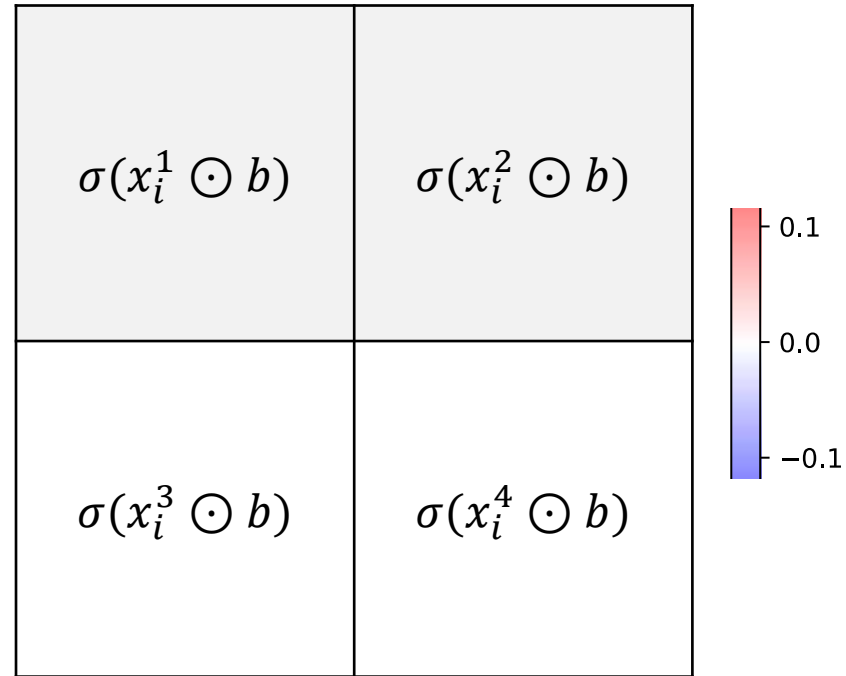







Examining filter output  $\sigma(x_i^R \odot b)$ , where  $x_i^R$  is the portion of image  $i$  where the filter is placed.





Examining filter output  $\sigma(x_i^R \odot b)$ , where  $x_i^R$  is the portion of image  $i$  where the filter is placed.





Examining filter output  $\sigma(x_i^R \odot b)$ , where  $x_i^R$  is the portion of image  $i$  where the filter is placed.



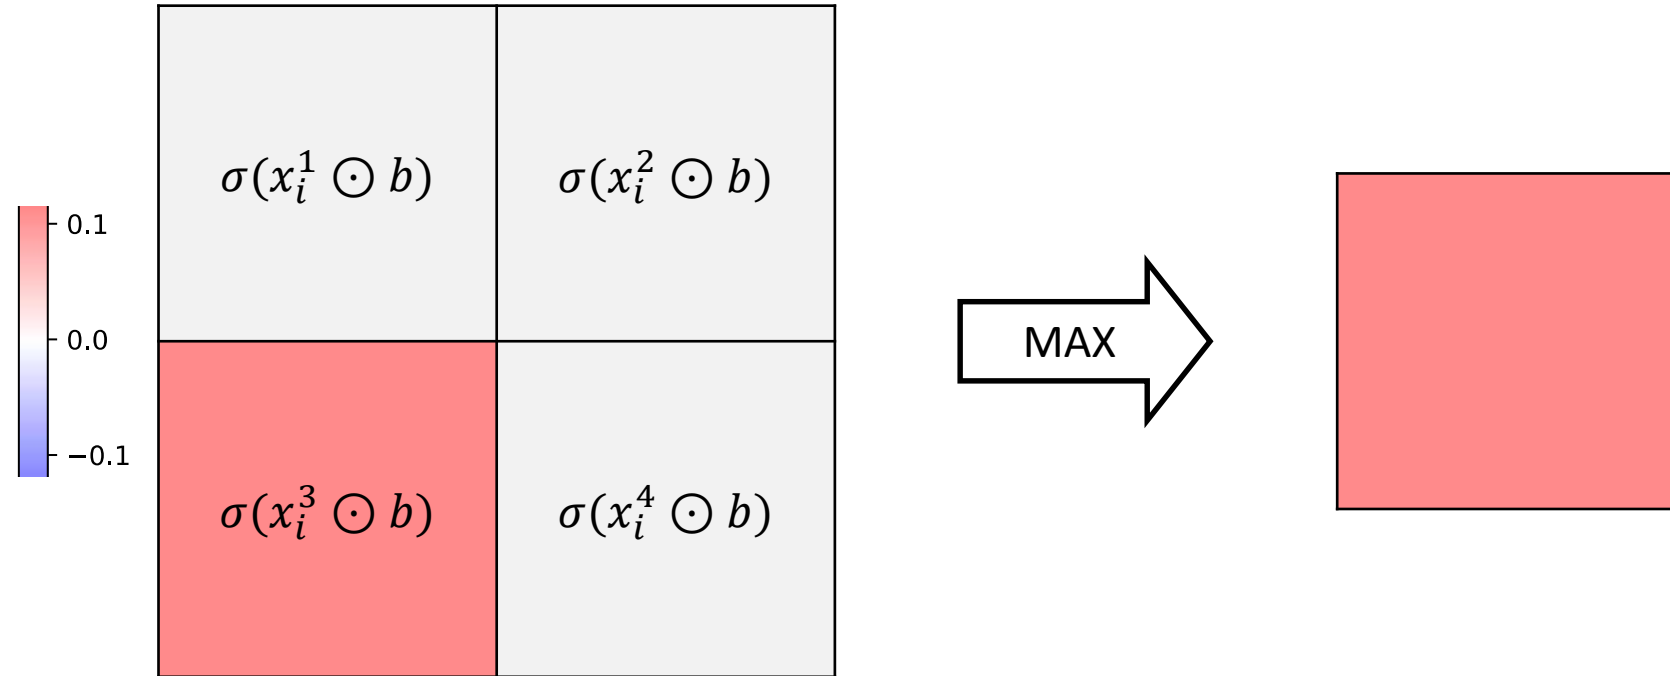
Examining filter output  $\sigma(x_i^R \odot b)$ , where  $x_i^R$  is the portion of image  $i$  where the filter is placed.



What if we want to know if a 1 is present  
*anywhere* in the image?

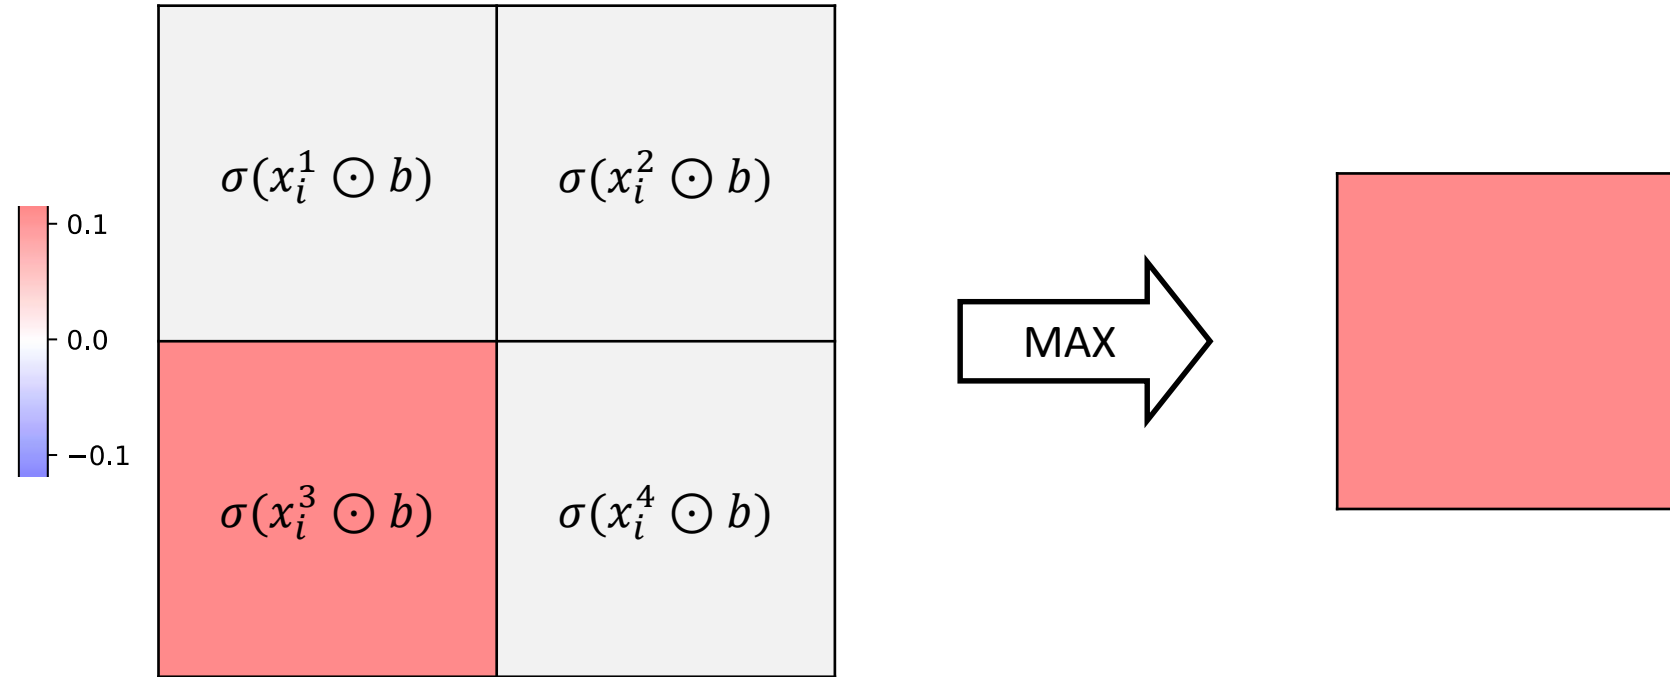


What if we want to know if a 1 is present  
*anywhere* in the image?

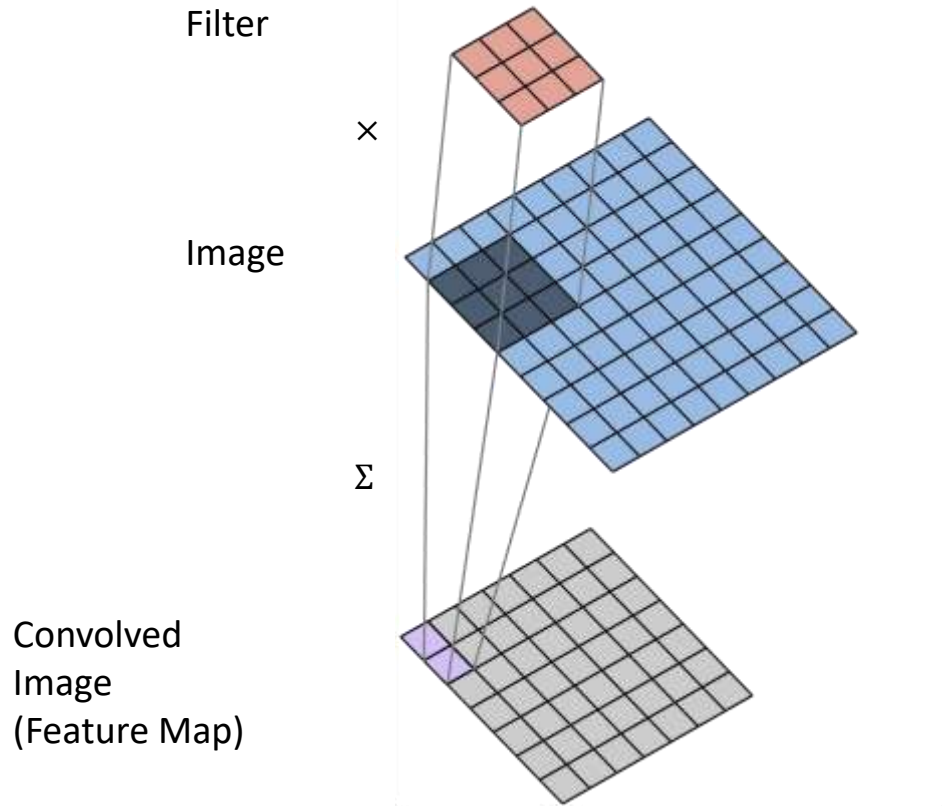




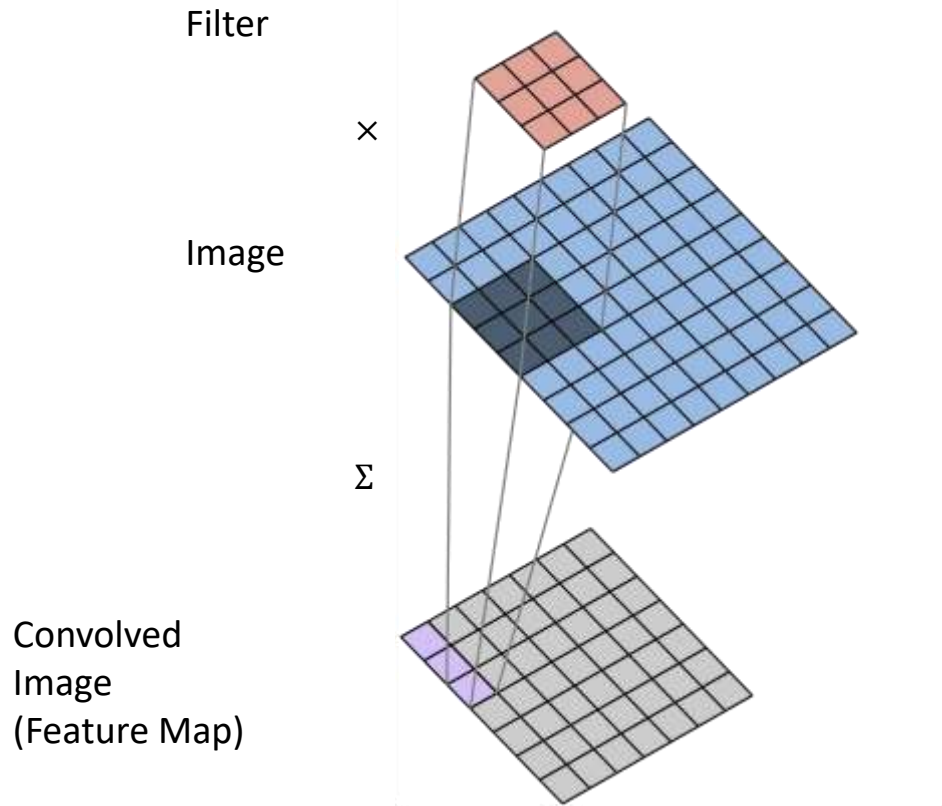
This idea is called “max pooling”, and is widely used in CNNs to determine whether features are present in a given region



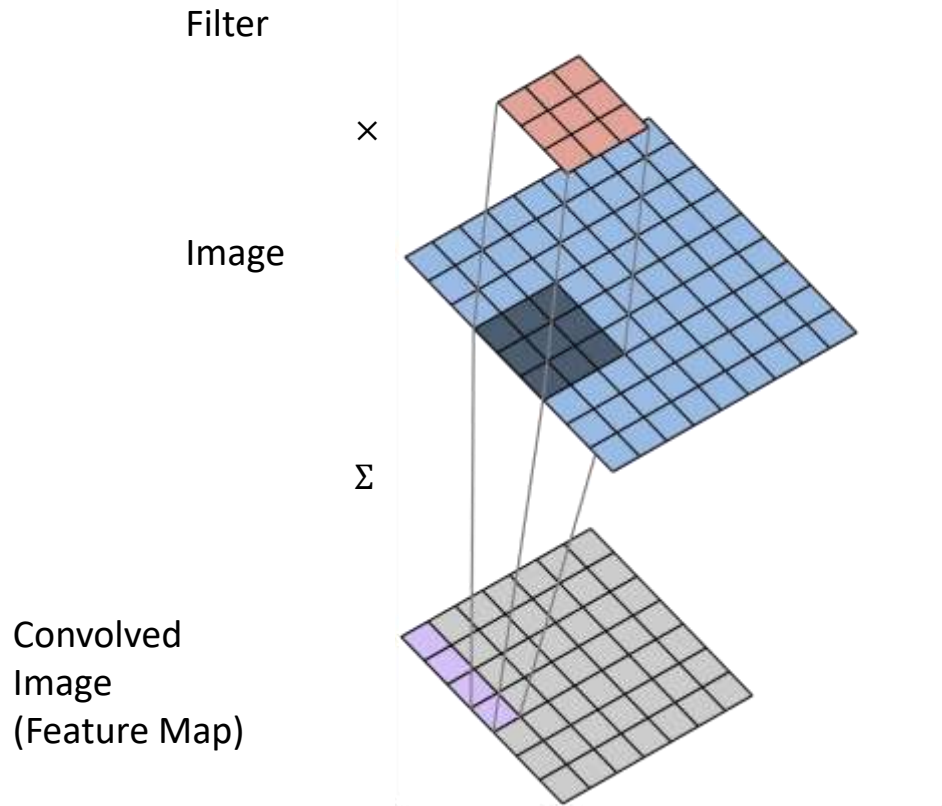
We perform “2D Spatial Convolution”  
as we move the filter across the image.



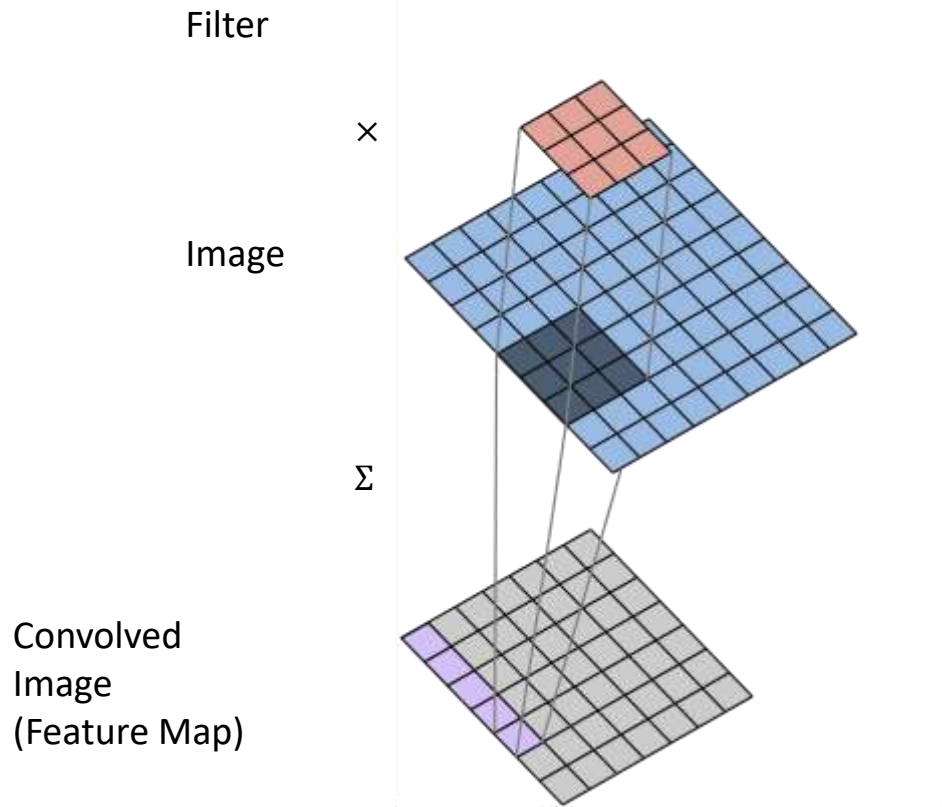
We perform “2D Spatial Convolution”  
as we move the filter across the image.




We perform “2D Spatial Convolution”  
as we move the filter across the image.




We perform “2D Spatial Convolution”  
as we move the filter across the image.



# An Example...



filter



image

# An Example...


|    |    |    |
|----|----|----|
| -1 | 1  | -1 |
| 1  | -1 | 1  |
| -1 | 1  | -1 |

filter


|    |    |    |    |    |
|----|----|----|----|----|
| -1 | -1 | -1 | -1 | -1 |
| -1 | -1 | -1 | 1  | -1 |
| -1 | -1 | 1  | -1 | 1  |
| -1 | -1 | -1 | 1  | -1 |
| -1 | -1 | -1 | -1 | -1 |

image

# An Example...



filter




image

|    |  |  |
|----|--|--|
| -1 |  |  |
|    |  |  |
|    |  |  |


$$x_i^R \odot b$$



# An Example...



filter




image


|    |   |  |
|----|---|--|
| -1 | 5 |  |
|    |   |  |
|    |   |  |

$$x_i^R \odot b$$

# An Example...



filter




image


|    |   |    |
|----|---|----|
| -1 | 5 | -5 |
|    |   |    |
|    |   |    |

$$x_i^R \odot b$$

# An Example...



filter




image


|    |   |    |
|----|---|----|
| -1 | 5 | -5 |
| 3  |   |    |
|    |   |    |

$$x_i^R \odot b$$

# An Example...



filter




image


|    |    |    |
|----|----|----|
| -1 | 5  | -5 |
| 3  | -5 |    |
|    |    |    |

$$x_i^R \odot b$$

# An Example...



filter




image


|    |    |    |
|----|----|----|
| -1 | 5  | -5 |
| 3  | -5 | 9  |
|    |    |    |

$$x_i^R \odot b$$

# An Example...



filter




image


|    |    |    |
|----|----|----|
| -1 | 5  | -5 |
| 3  | -5 | 9  |
| -1 |    |    |

$$x_i^R \odot b$$

# An Example...



filter




image


|    |    |    |
|----|----|----|
| -1 | 5  | -5 |
| 3  | -5 | 9  |
| -1 | 5  |    |

$$x_i^R \odot b$$

# An Example...



filter




image

|    |    |    |
|----|----|----|
| -1 | 5  | -5 |
| 3  | -5 | 9  |
| -1 | 5  | -5 |


$$x_i^R \odot b$$



# An Example...



filter




image


|    |    |    |
|----|----|----|
| -1 | 5  | -5 |
| 3  | -5 | 9  |
| -1 | 5  | -5 |

$$x_i^R \odot b$$


Each location where the filter was centered has been evaluated: “how similar is this location to the filter”?



filter




image





$x_i^R \odot b$


# Convolutional Filters Are Feature Detectors





- Now we know how to identify a “1” or a “0” anywhere in an image.
- What if we want to identify a “10”?
- Option 1: Design a new filter for “10”
- Option 2: Utilize our “1” and “0” filters...





"1" filter output



"0" filter output



Our "10" filter looks for a match from the "1" filter to the left of a match from the "0" filter





**In this way, we learn to identify a hierarchy of features  
rather than a huge number of complex features**

# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images







# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images





# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images





**Low-level structure:** lines,  
curves

# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images





**Low-level structure:** lines,  
curves

# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images





**Mid-level structure:** shapes

# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images



**Mid-level structure:** shapes


# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images




**High-level structure:** groups of shapes

# CNNs Take Advantage of **Repeated, Hierarchical Structure** in Images

**Bear**






**Train**




**High-level structure:** groups of shapes → objects

Consider a Set of “Toy” Images,  
for illustration of how this structure can be extracted by an algorithm









...















# High-Level Motifs/Structure










# Shared Substructure Within Motifs




# Shared Substructure Within Motifs



# Shared Substructure Within Motifs




# Shared Substructure Within Motifs




# Hierarchical Representation of Images

Layer 3:  
Motifs






Layer 2:  
Sub-Motifs




Layer 1:  
Fundamental Building Blocks




# Recall the Data/Images



...




# Convolutional Filter







# Convolutional Filter





# Convolutional Filter





# Multiple Filters, One for Each Building Block









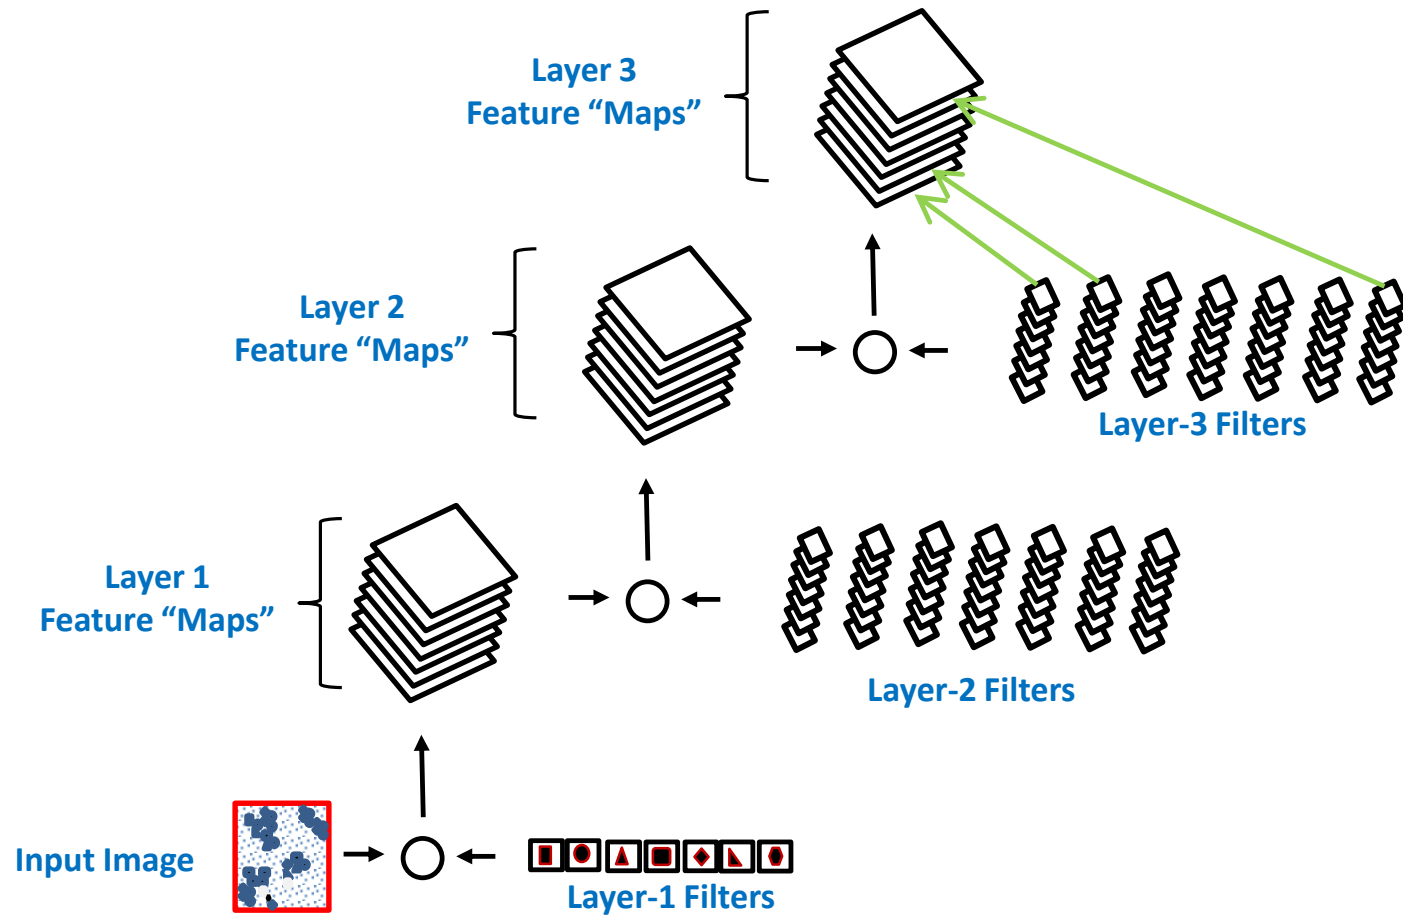














# Deep CNN Architecture

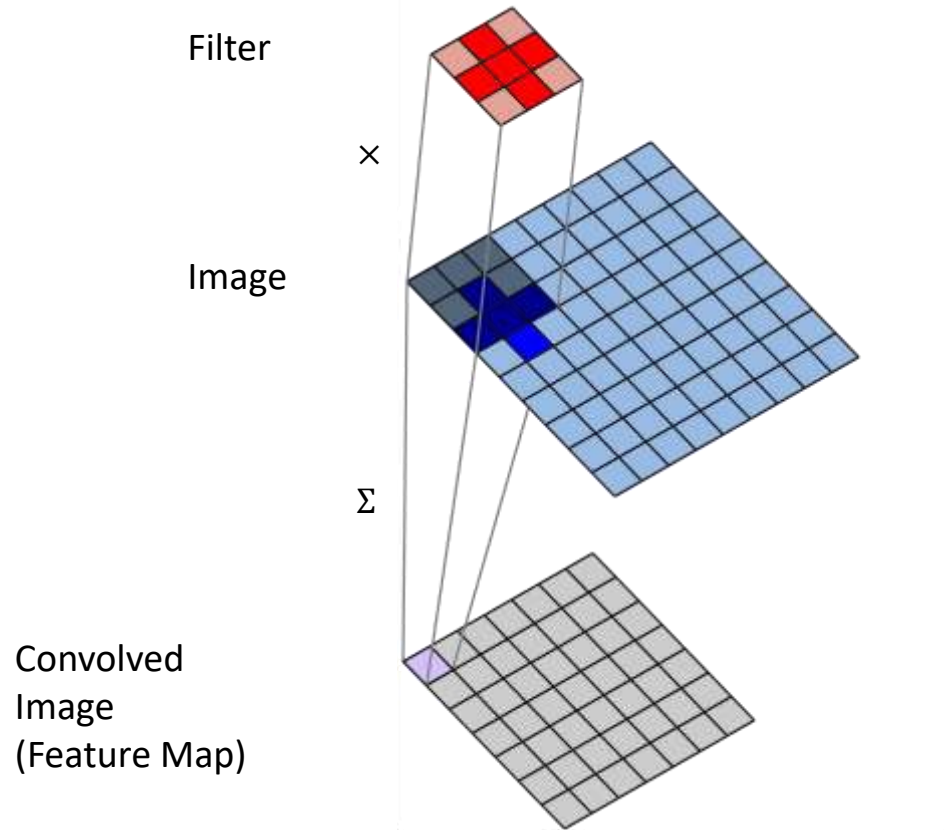


# Advantage of Hierarchical Features?




- By learning and sharing statistical similarities within high-level motifs, we better leverage all training data
- If we do not use such a hierarchy, top-level motifs would be learned in isolation of each other

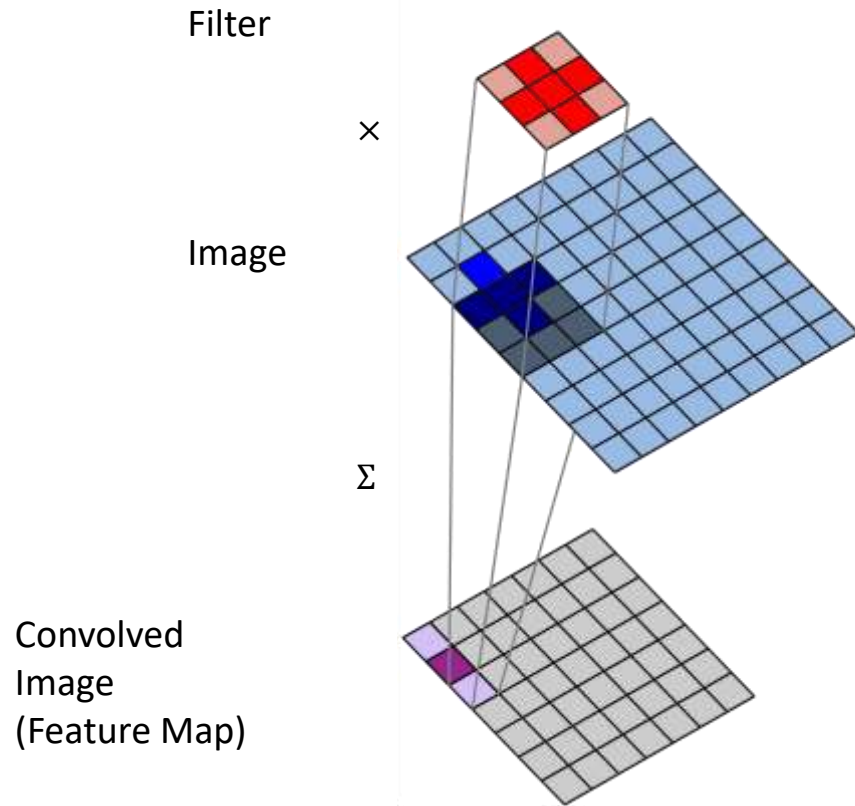
## 2D Spatial Convolution



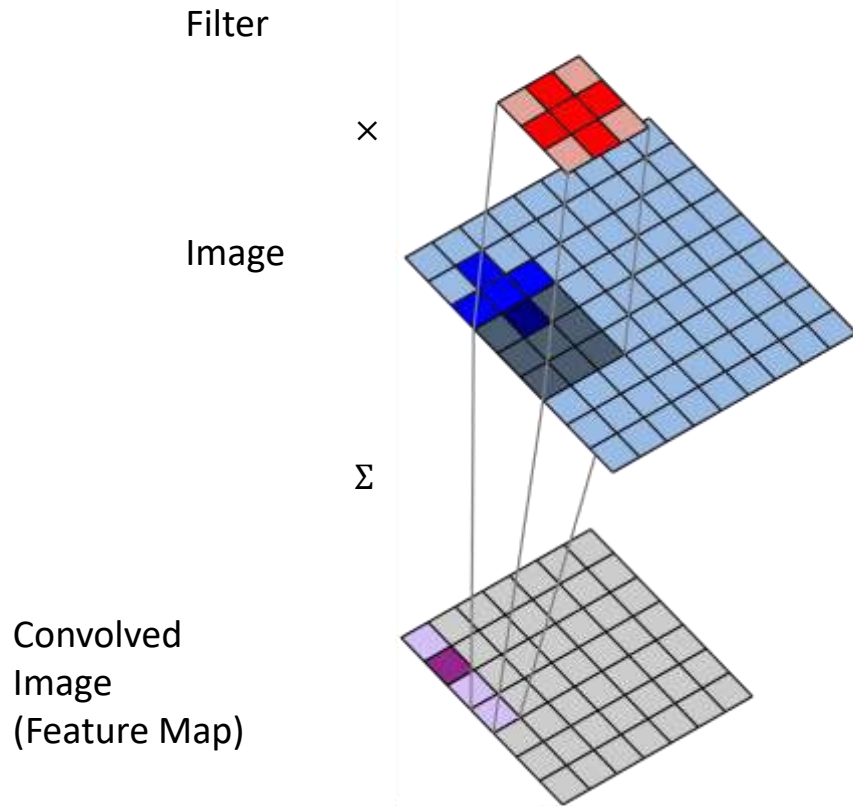
## 2D Spatial Convolution



# 2D Spatial Convolution

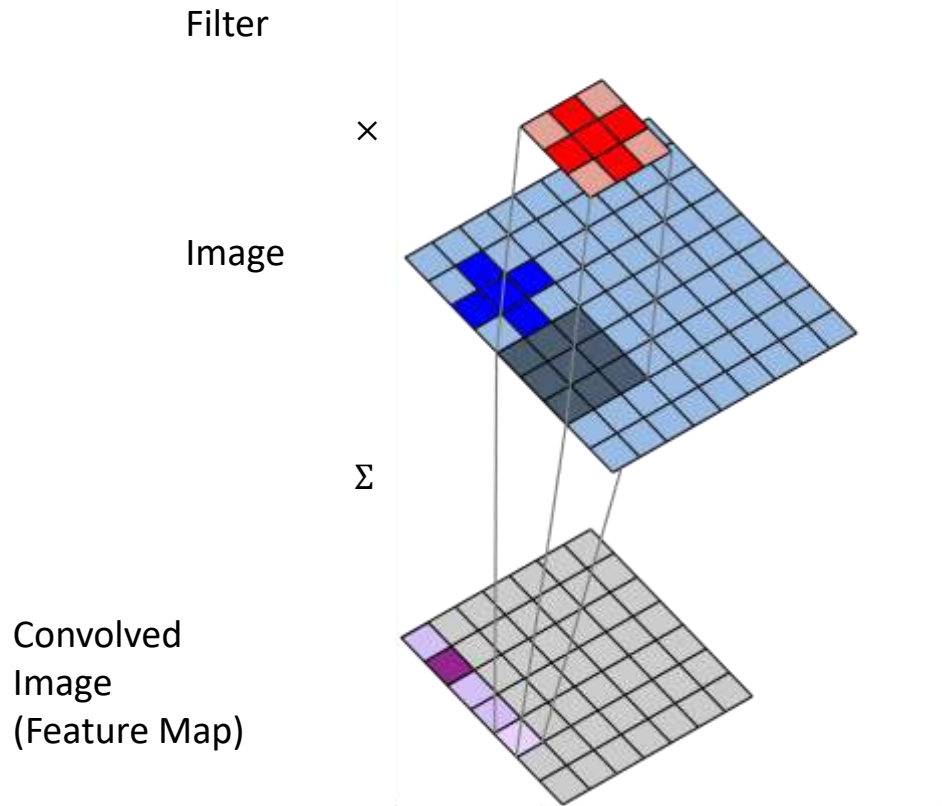


# 2D Spatial Convolution





# 2D Spatial Convolution



# 2D Spatial Convolution

Filter ( $W$ )

|   |    |   |
|---|----|---|
| 0 | 1  | 0 |
| 1 | -4 | 1 |
| 0 | 1  | 0 |

Input ( $I$ )

|   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|
| 0 | 1 | 0 | 0 | 1 | 1 | 4 | 1 | 1 |
| 9 | 2 | 1 | 1 | 8 | 8 | 2 | 2 | 6 |
| 2 | 1 | 4 | 0 | 1 | 0 | 0 | 1 | 5 |
| 0 | 1 | 5 | 9 | 1 | 2 | 7 | 1 | 4 |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| 0 | 1 | 0 | 0 | 1 | 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 0 | 1 | 0 | 0 | 1 | 0 |
| 1 | 8 | 1 | 1 | 8 | 1 | 1 | 9 | 1 |
| 0 | 1 | 0 | 0 | 1 | 0 | 0 | 1 | 0 |



|     |      |     |   |   |   |   |   |   |
|-----|------|-----|---|---|---|---|---|---|
| 0*0 | 1*1  | 0*0 | 0 | 1 | 1 | 4 | 1 | 1 |
| 1*9 | -4*2 | 1*1 | 1 | 8 | 8 | 2 | 2 | 6 |
| 2*0 | 1*1  | 4*0 | 0 | 1 | 0 | 0 | 1 | 5 |
| 0   | 1    | 5   | 9 | 1 | 2 | 7 | 1 | 4 |
| 1   | 1    | 1   | 1 | 1 | 1 | 1 | 1 | 1 |
| 0   | 1    | 0   | 0 | 1 | 0 | 0 | 1 | 0 |
| 0   | 1    | 0   | 0 | 1 | 0 | 0 | 1 | 0 |
| 1   | 8    | 1   | 1 | 8 | 1 | 1 | 9 | 1 |
| 0   | 1    | 0   | 0 | 1 | 0 | 0 | 1 | 0 |

$$\sum_i \sum_j I_{ij} W_{ij}$$

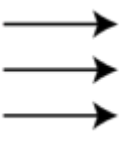


|   |  |  |  |  |  |  |  |  |
|---|--|--|--|--|--|--|--|--|
| 4 |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |
|   |  |  |  |  |  |  |  |  |


# 2D Spatial Convolution

Filter ( $W$ )


|   |    |   |
|---|----|---|
| 0 | 1  | 0 |
| 1 | -4 | 1 |
| 0 | 1  | 0 |




Input ( $I$ )




# Filters Operate Over Input Volumes




# Filters Operate Over Input Volumes




# Filters Operate Over Input Volumes



# Given Labeled Training Images, How do we Learn the Parameters of the CNN?




Training Set



$$p_i = \sigma(b_0 + b_1x_{i1} + b_2x_{i2} + \dots + b_Mx_{iM})$$

Untrained Logistic Regression  
Model (or "Network")



$$b = (b_0, \dots, b_M)$$

Trained Model (with  
learned parameters)

# Given Labeled Training Images, How do we Learn the Parameters of the CNN?



CNN  
Architecture



Trained  
CNN



# Architecture (specified) vs Parameters (learned)

## Architecture:

- Number of layers
- Layer types (e.g. convolutional, pooling, fully connected)
- Number of filters in each layer
- Shape and size of filters

Use 3x3 filters  
In layer 1


|  |  |  |
|--|--|--|
|  |  |  |
|  |  |  |
|  |  |  |

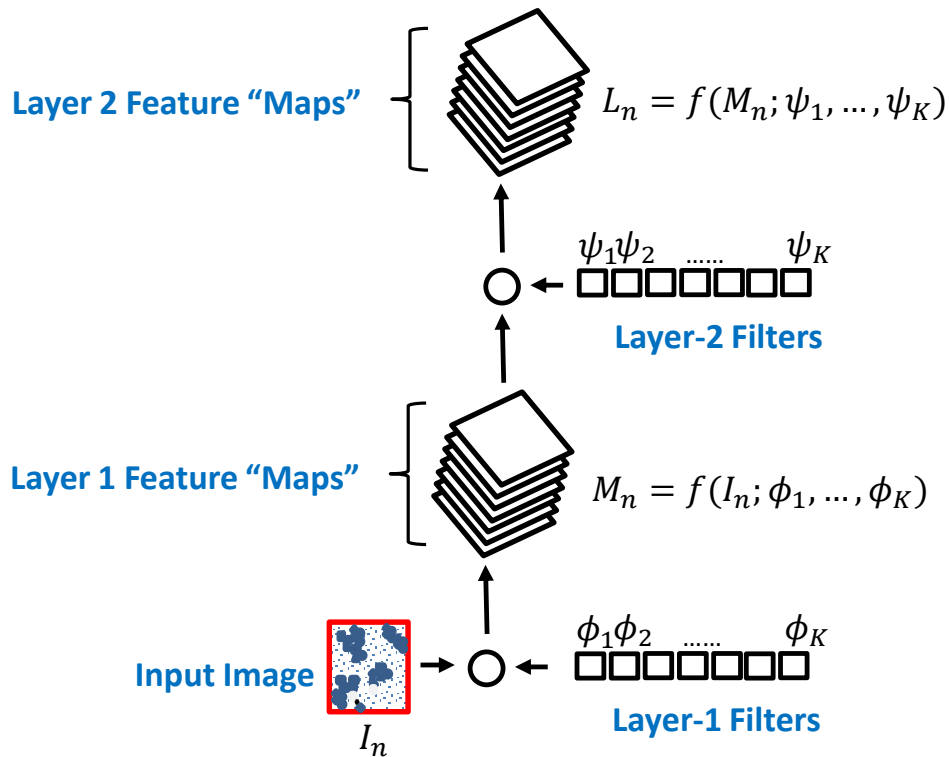
## Parameters:

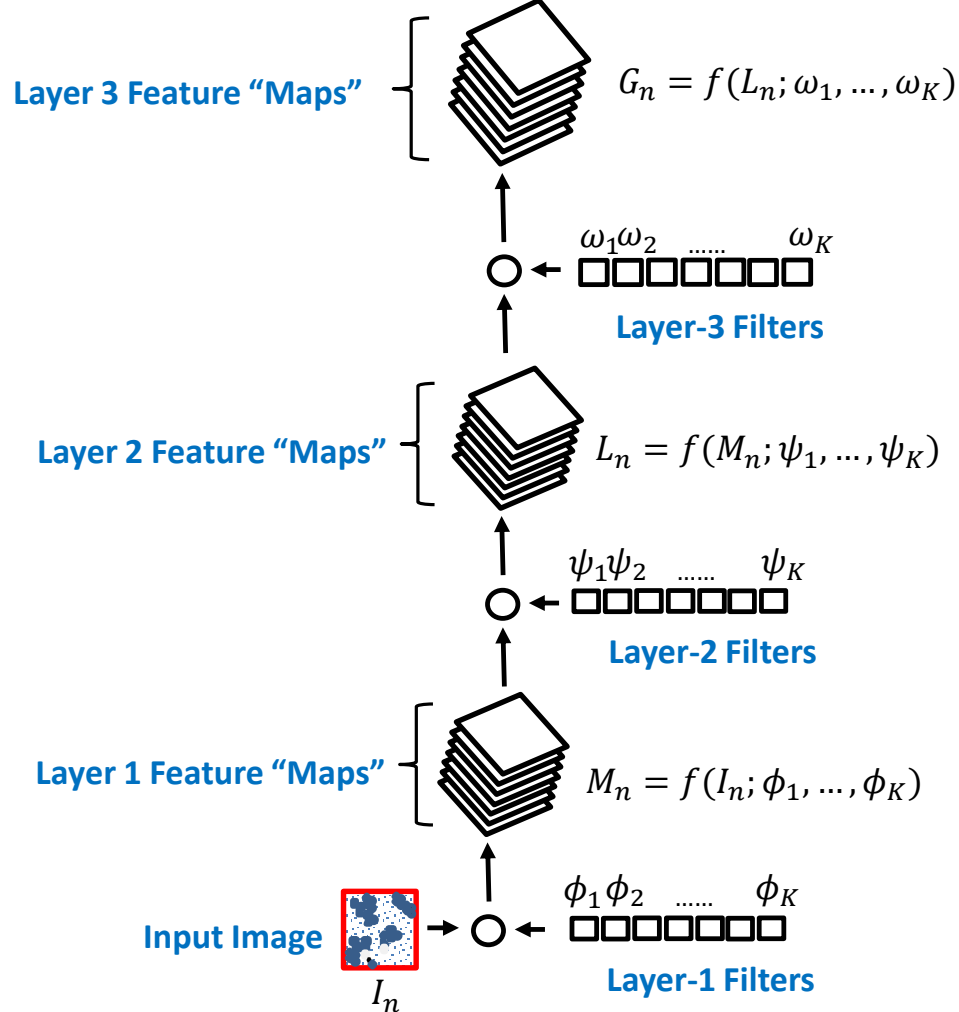
- Individual Elements of each filter
- Parameters of other layers

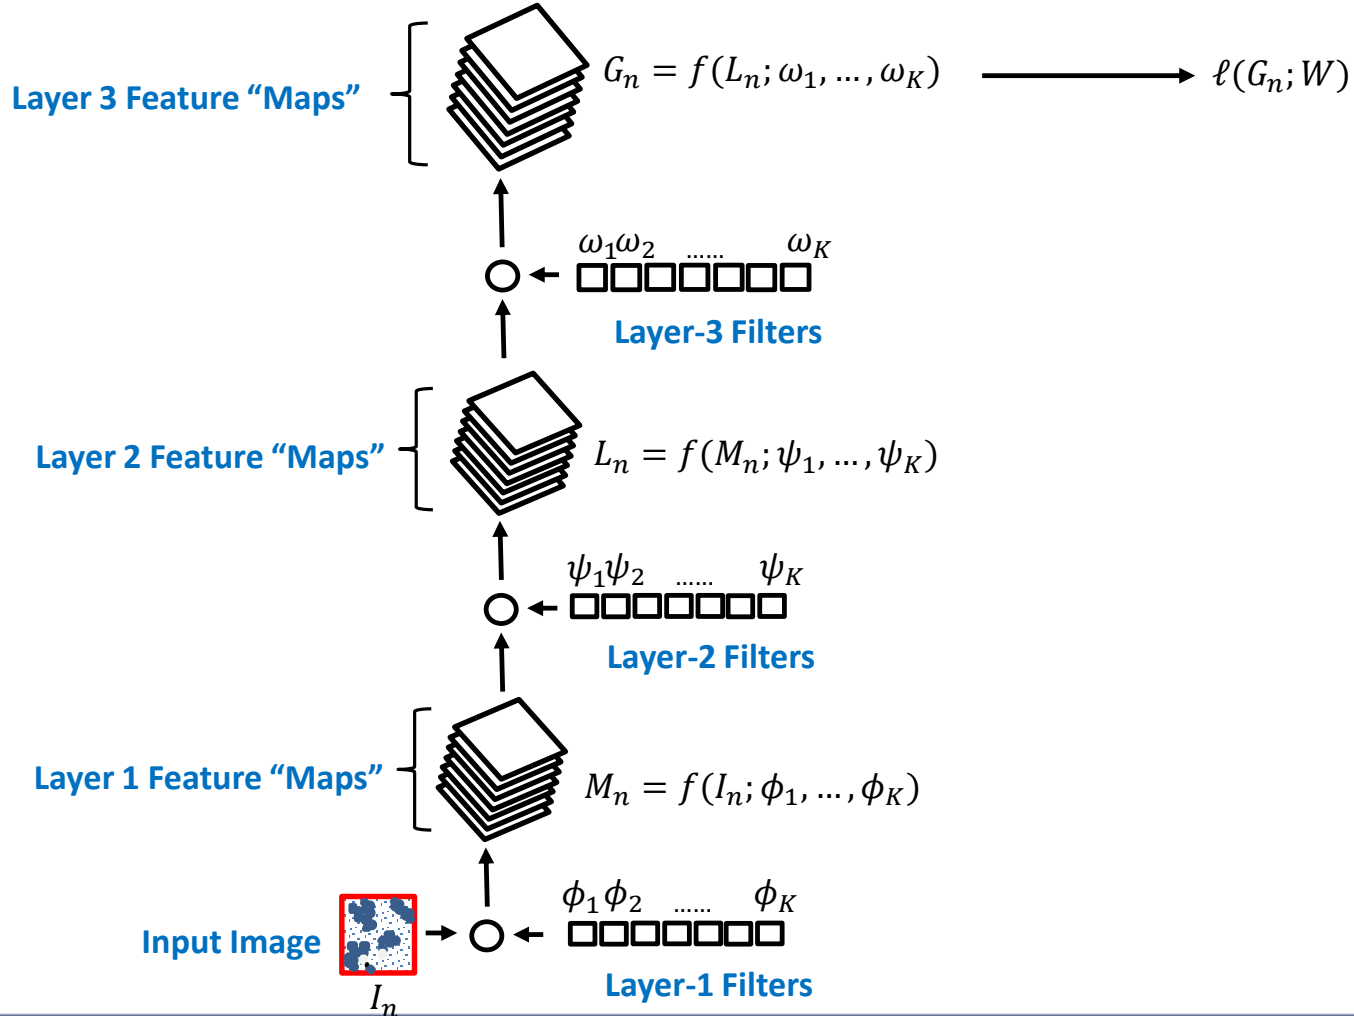
Learn values of  
Each layer 1 filter

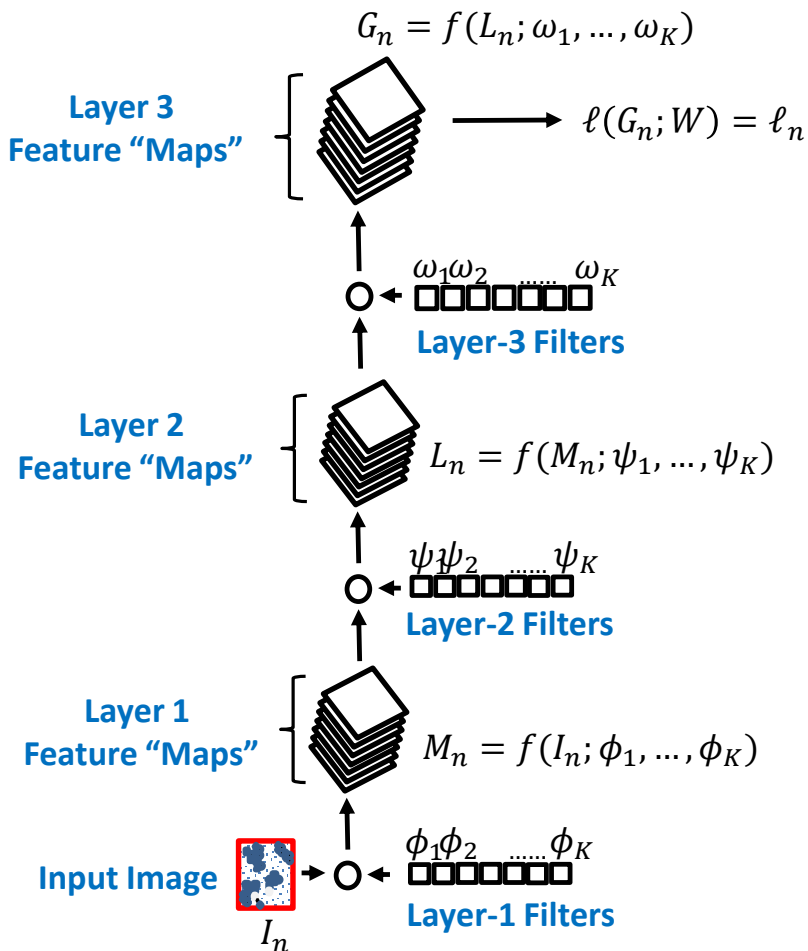
|    |    |    |
|----|----|----|
| -1 | 1  | -1 |
| 1  | -1 | 1  |
| -1 | 1  | -1 |











- Assume we have labeled images  $\{I_n, y_n\}_{n=1, N}$
- $I_n$  is image  $n$ ,  $y_n \in \{+1, -1\}$  is associated label
- Average loss, which depends on model parameters:

$$\frac{1}{N} \sum_{n=1}^N \text{loss}(y_n, \ell_n)$$

- Find specific parameters that minimize the average loss

# Summary

- Convolutional neural networks learn to recognize **high-level structure** in images by building **hierarchical representations of features**
- Features are extracted via spatial convolutions with **filters**
- Filters are learned via iterative minimization of a loss function
- Convolutional neural networks have shown capabilities beyond human performance for image analysis