

A brief note on model learning

Matthew Engelhard

How does training, or learning, work?

I want to find parameters that give me the *best* performance.

1. What does best mean, exactly? We need to be able to quantify it numerically in order to determine which parameters are best.
 - > define the *loss*
 - > the *best* parameters are those that minimize it
2. How do I find the parameters that minimize the loss?
 - > this is the field of *optimization*

Cross-Entropy Loss for Binary Classifier

Parameter Optimization by Gradient Descent